High-level concurrency concepts

Václav Pech

http://jroller.com/vaclav
http://www.vaclavpech.eu
@vaclav_pech

Why concurrency?

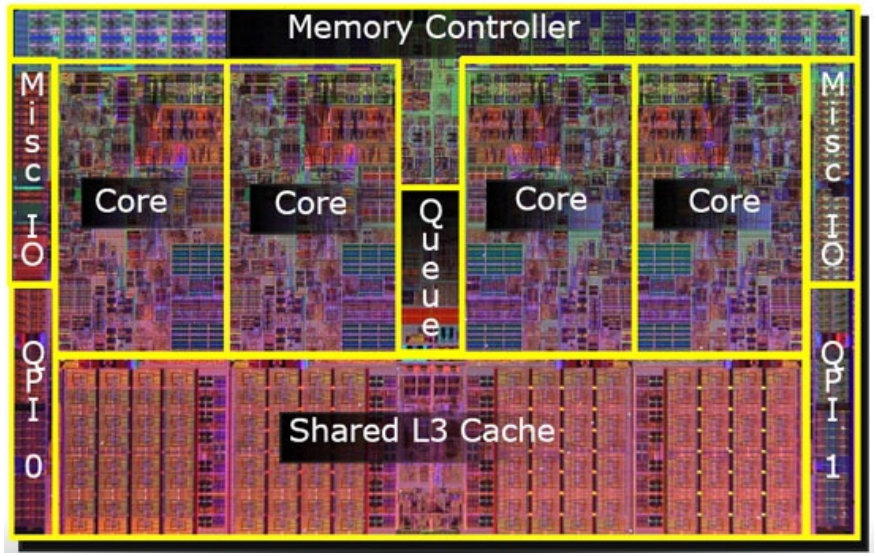

We’re all in the parallel computing business!

of cores

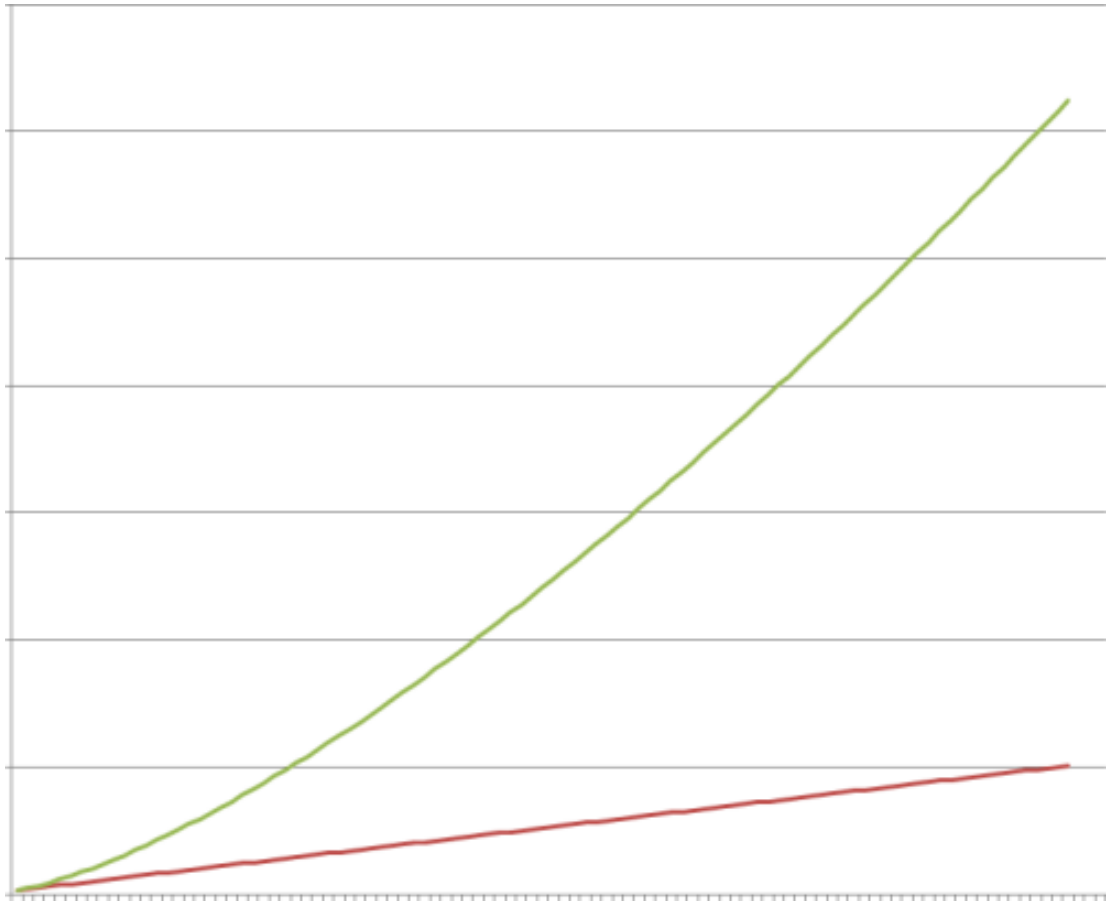

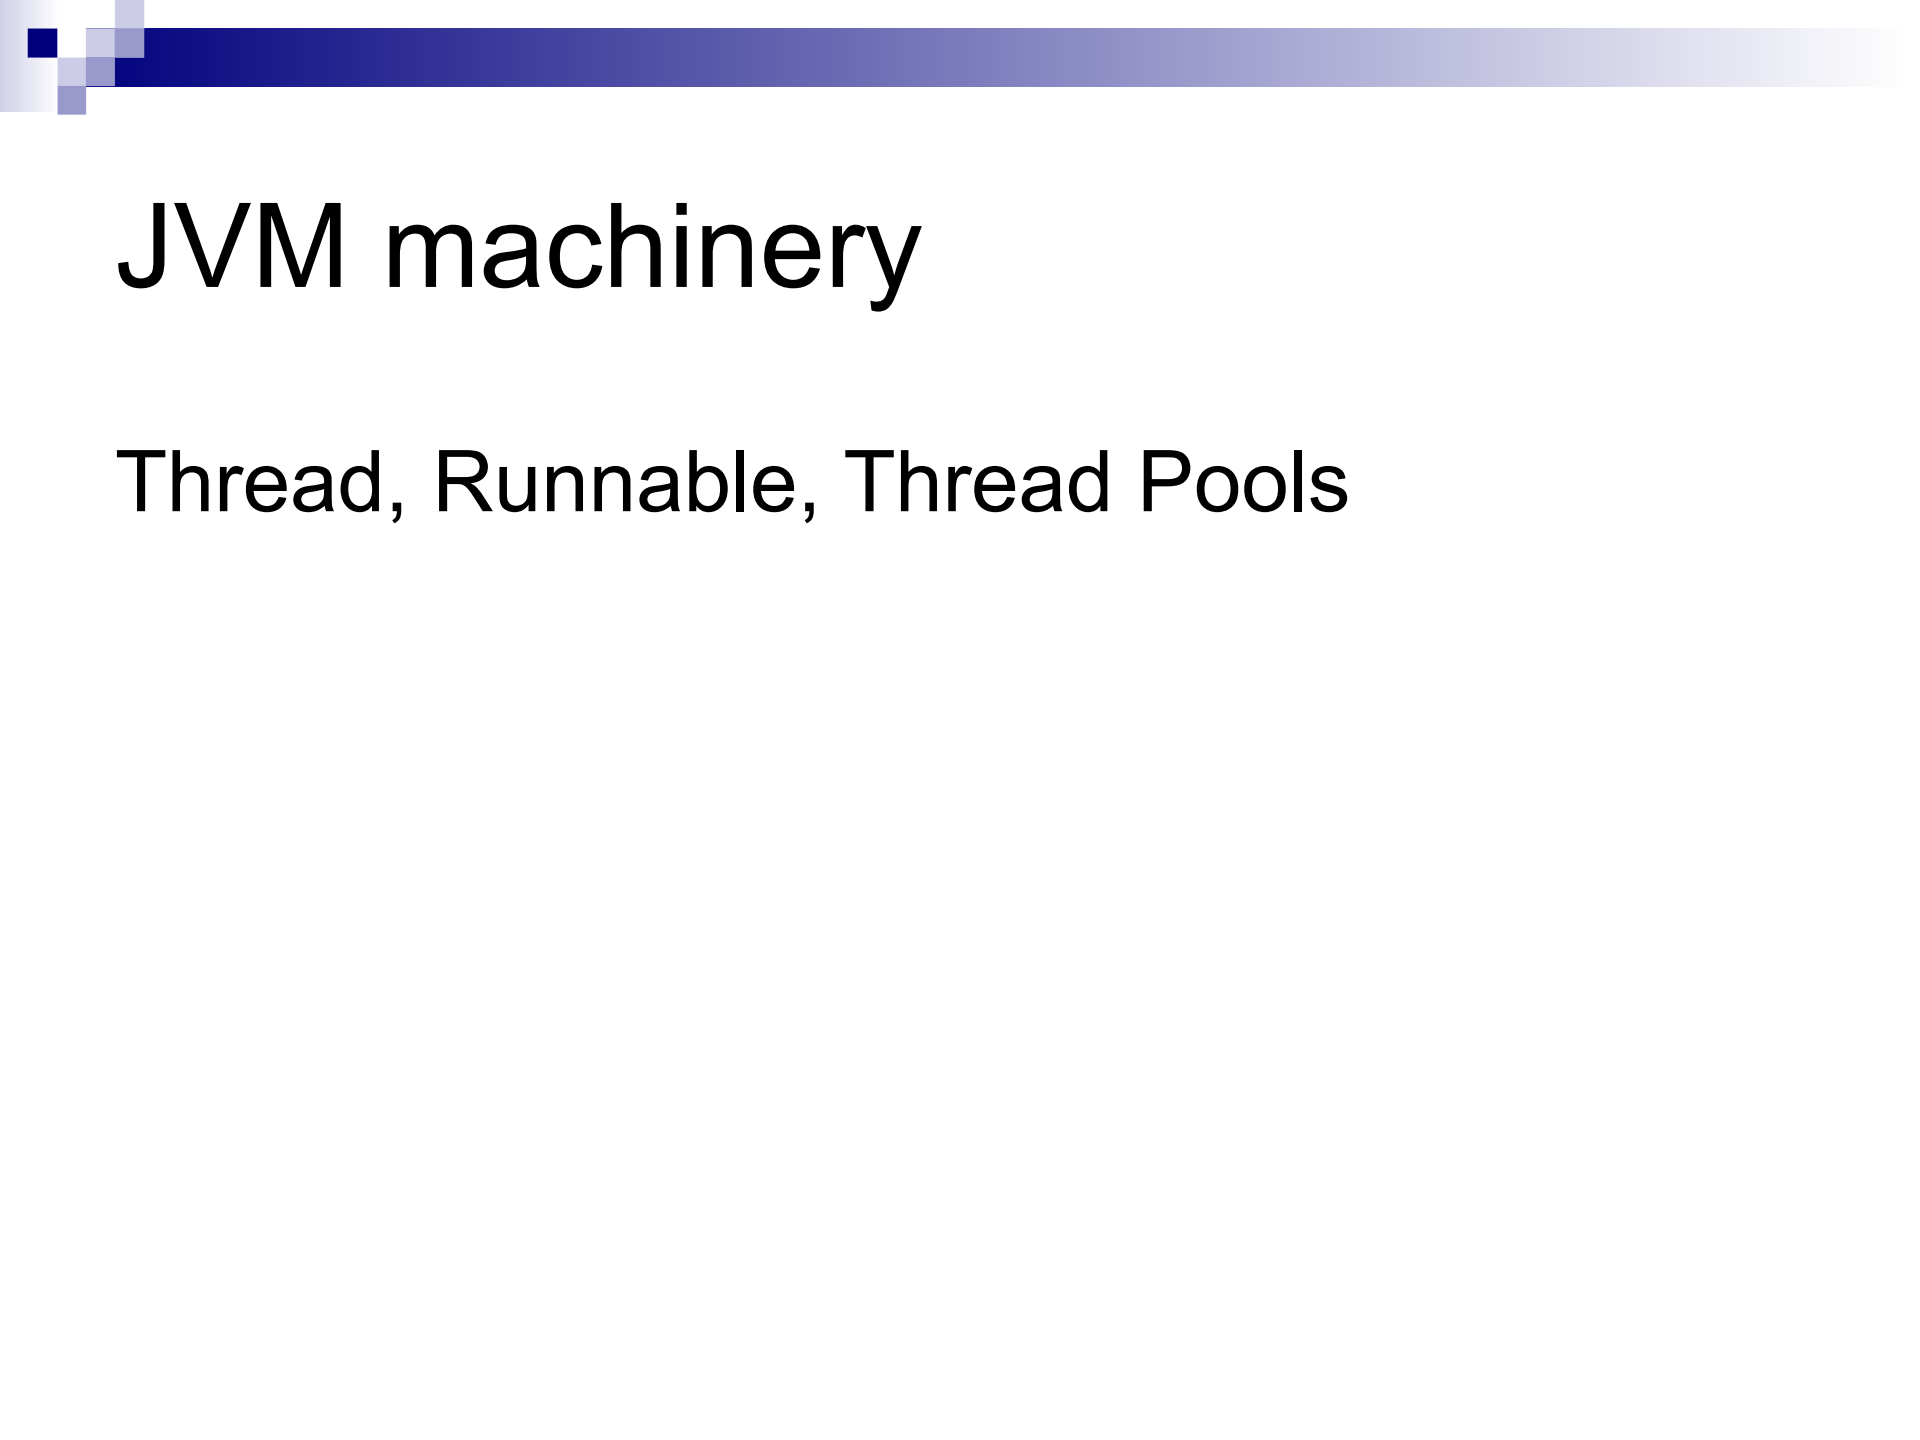
JVM machinery
Thread, Runnable, Thread Pools

JVM machinery

Thread, Runnable, Thread Pools

Synchronized blocks

Volatile

Locks

Atomic

Dealing with threads sucks!

```
public class Counter {
    private static long count = 0;

    public Counter() {

        count++;

    }
}
```

Dealing with threads sucks!

```
public class Counter {
    private volatile static long count = 0;

    public Counter() {

            count++;

    }
}
```

Dealing with threads sucks!

```
public class Counter {
    private volatile static long count = 0;

    public Counter() {

            count = count + 1;

    }
}
```

Dealing with threads sucks!

```
public class Counter {
    private static long count = 0;

    public Counter() {
        synchronized (this) {
            count++;
        }
    }
}
```

Dealing with threads sucks!

```
public class Counter {
    private static long count = 0;

    public Counter() {
        synchronized (this.getClass()) {
            count++;
        }
    }
}
```

Dealing with threads sucks!

```
public class Counter {
    private Long count = 0;

    public doSomething() {
        synchronized (count) {
            count++;
        }
    }
}
```

Dealing with threads sucks!

```
public class Counter {
    private Long count = 0;

    public doSomething() {
        synchronized (count) {
            count = new Long(count.longValue() + 1);
        }
    }
}
```

Dealing with threads sucks!

```
public class ClickCounter implements ActionListener {
    public ClickCounter(JButton button) {
        button.addActionListener(this);
    }

    public void actionPerformed(final ActionEvent e) {
        ...
    }
}
```

Stone age of parallel SW

Dead-locks

Live-locks

Race conditions

Starvation

Shared Mutable State

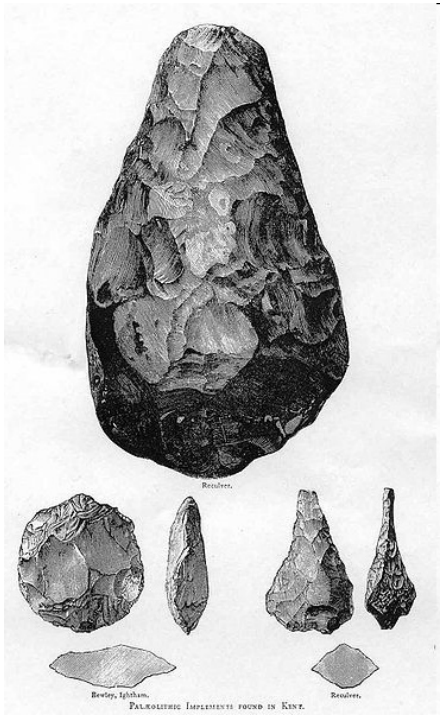

Why high-level concurrency?

Multithreaded programs today work mostly by accident!

Structure

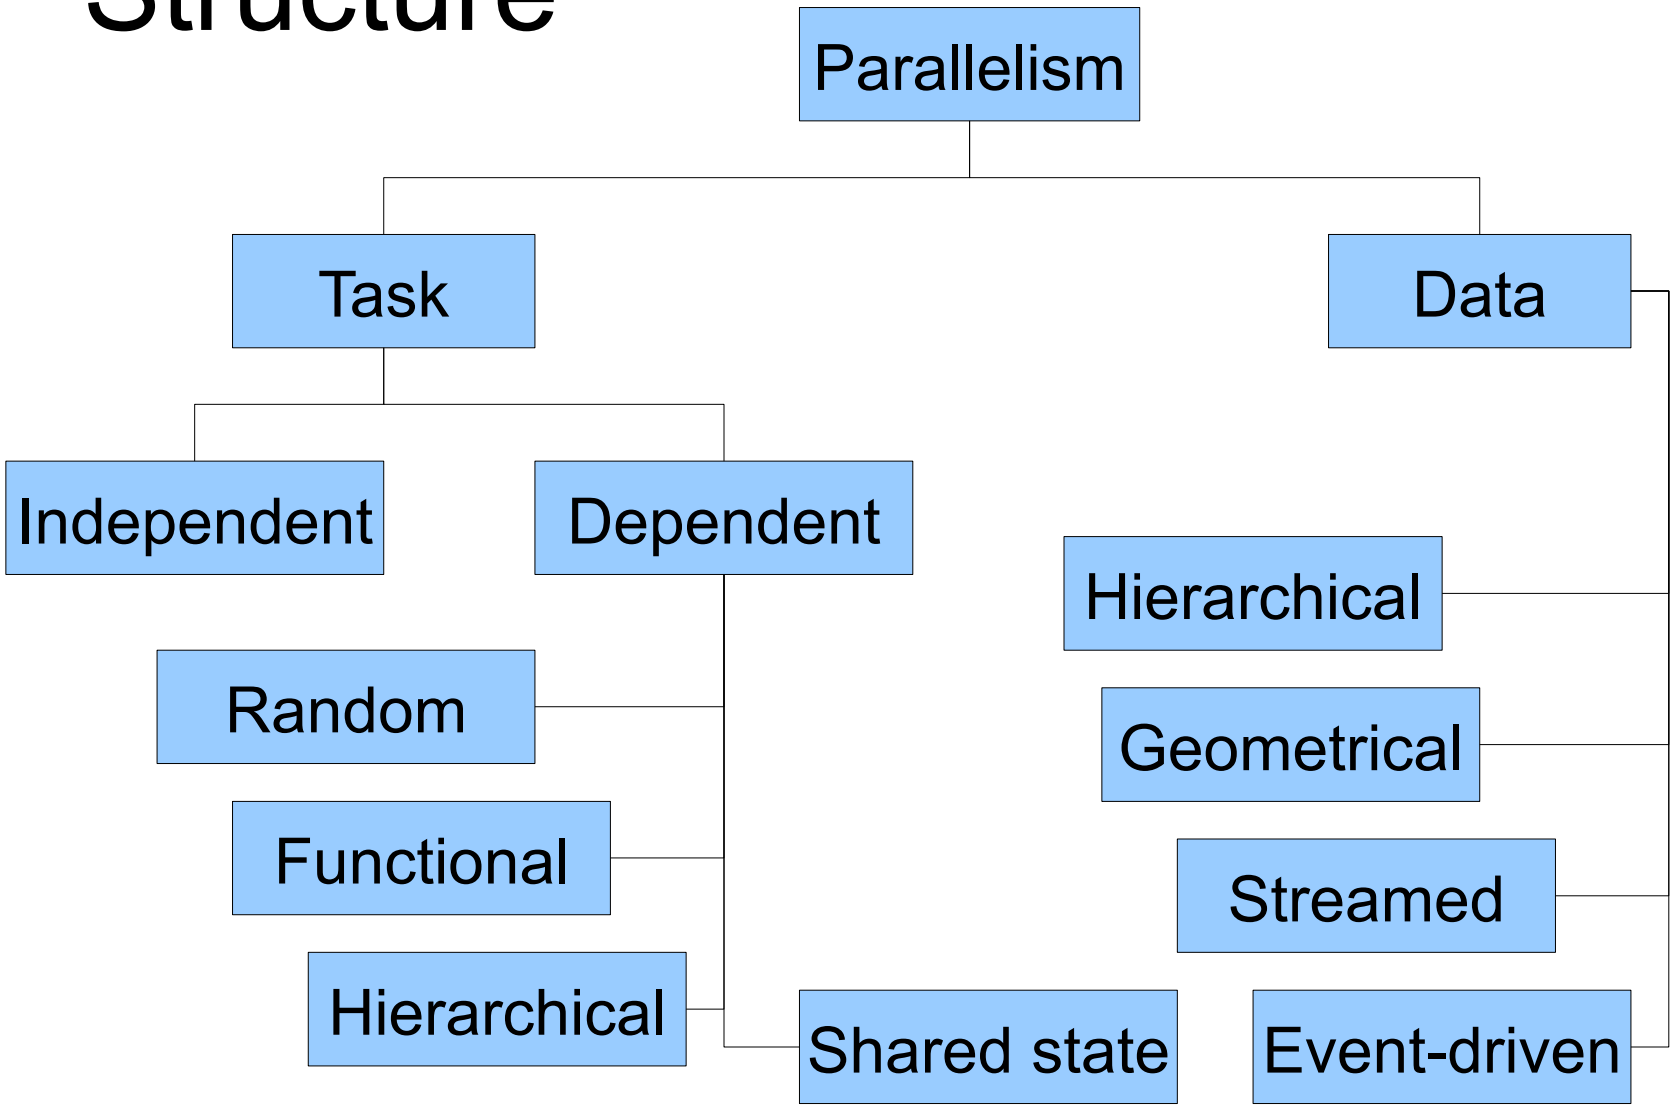
Parallelism
Task
Data
Independent
Dependent
Hierarchical
Random
Geometrical
Functional
Streamed
Hierarchical
Shared state
Event-driven

Structure

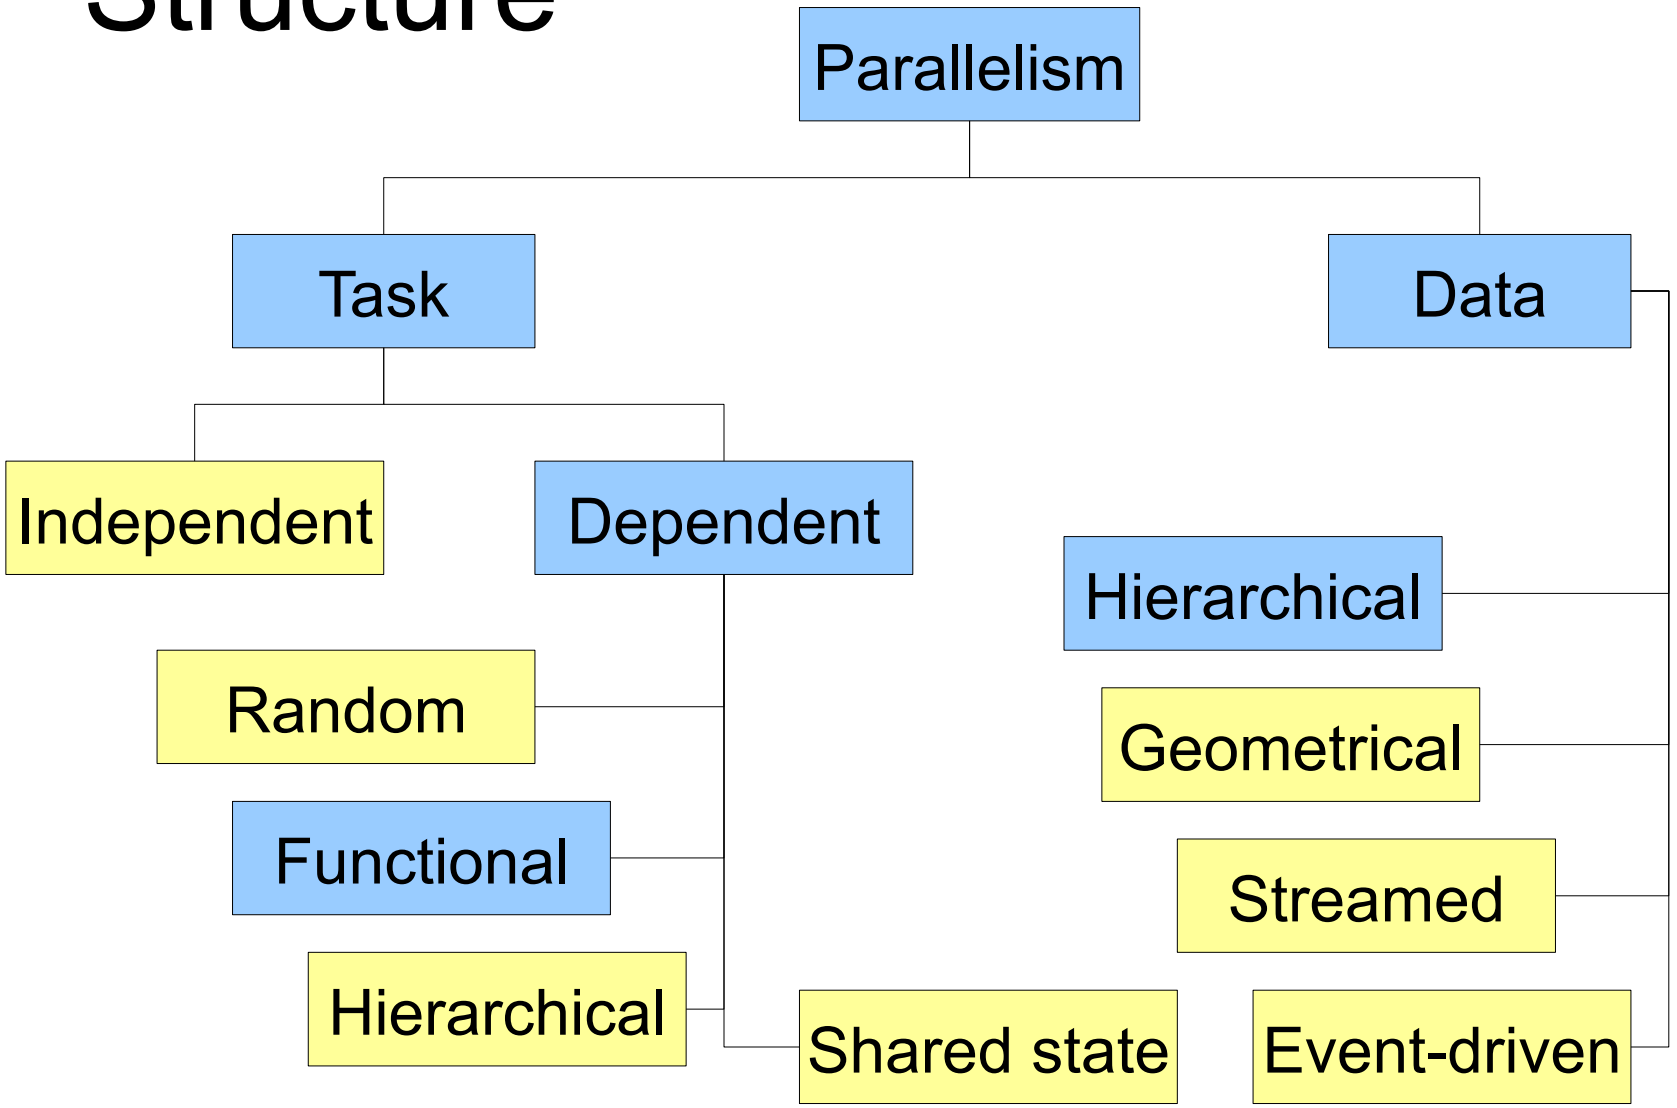
Parallelism
Task
Data
Independent
Dependent
Hierarchical
Random
Geometrical
Functional
Streamed
Hierarchical
Shared state
Event-driven

Task parallelism

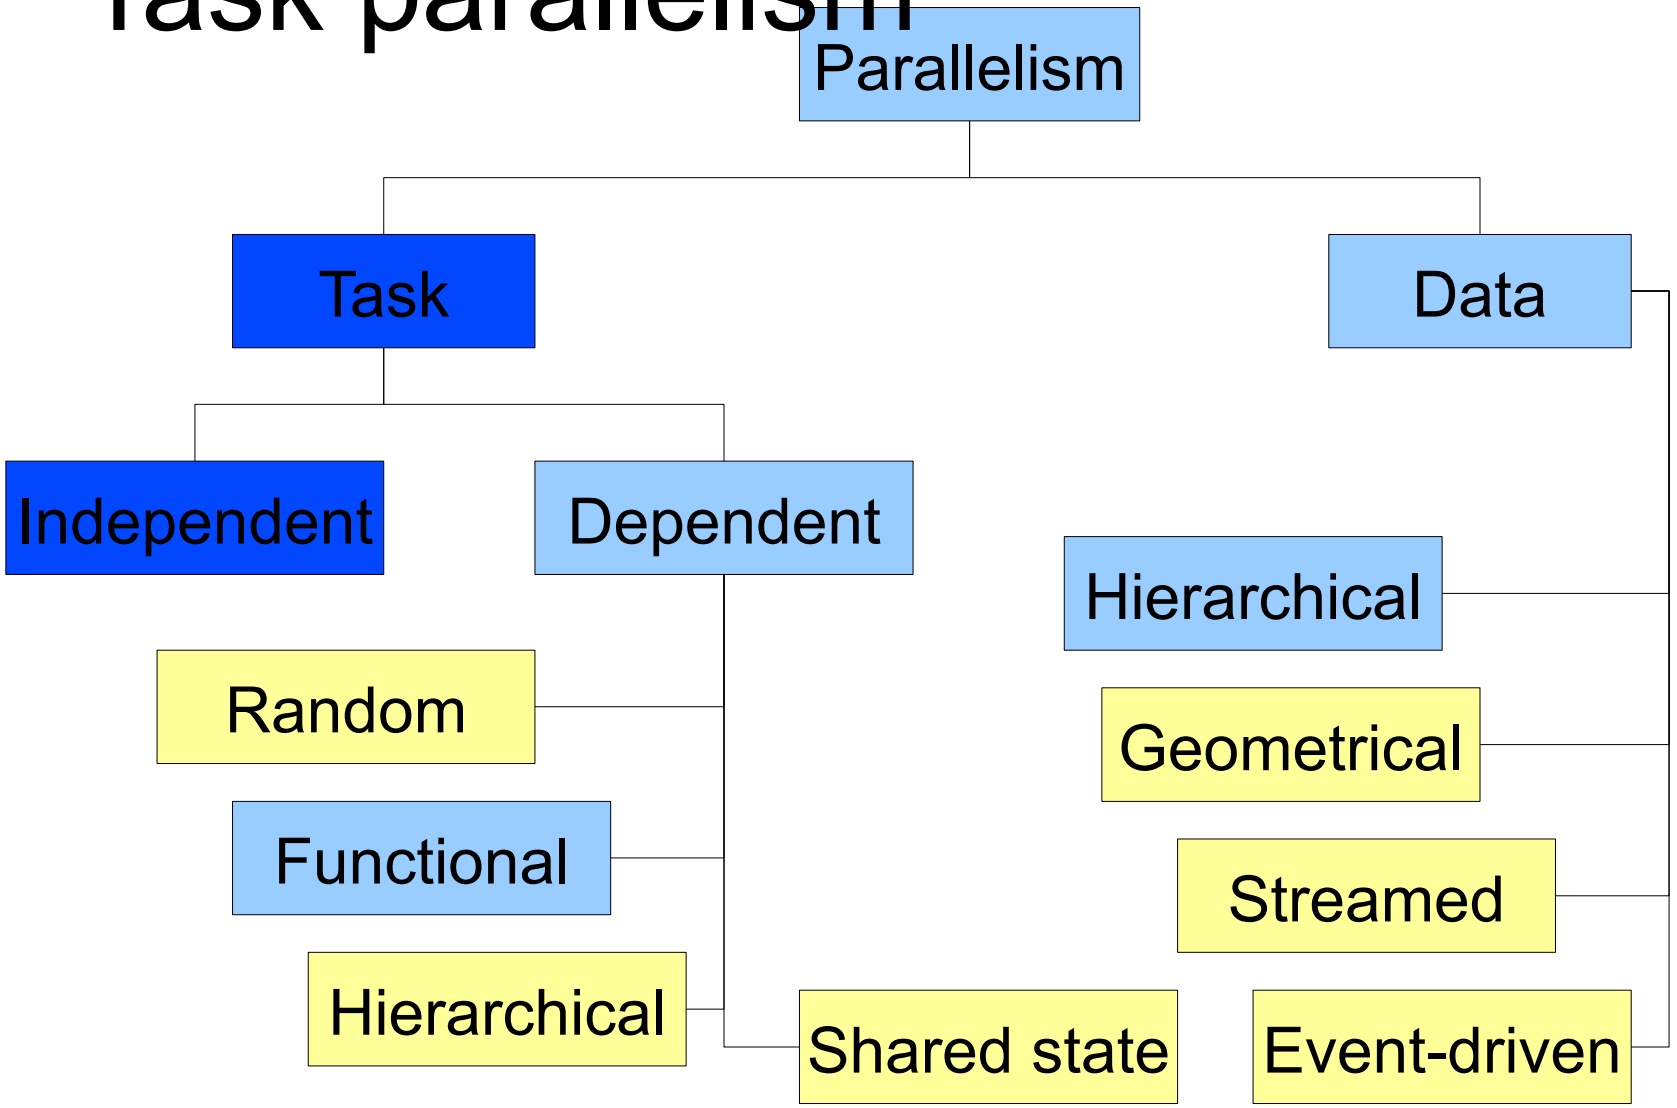

Asynchronous invocation

```
Future f = threadPool.submit(calculation);
…
System.out.println(“Result: “ + f.get());
```

Thread Pool

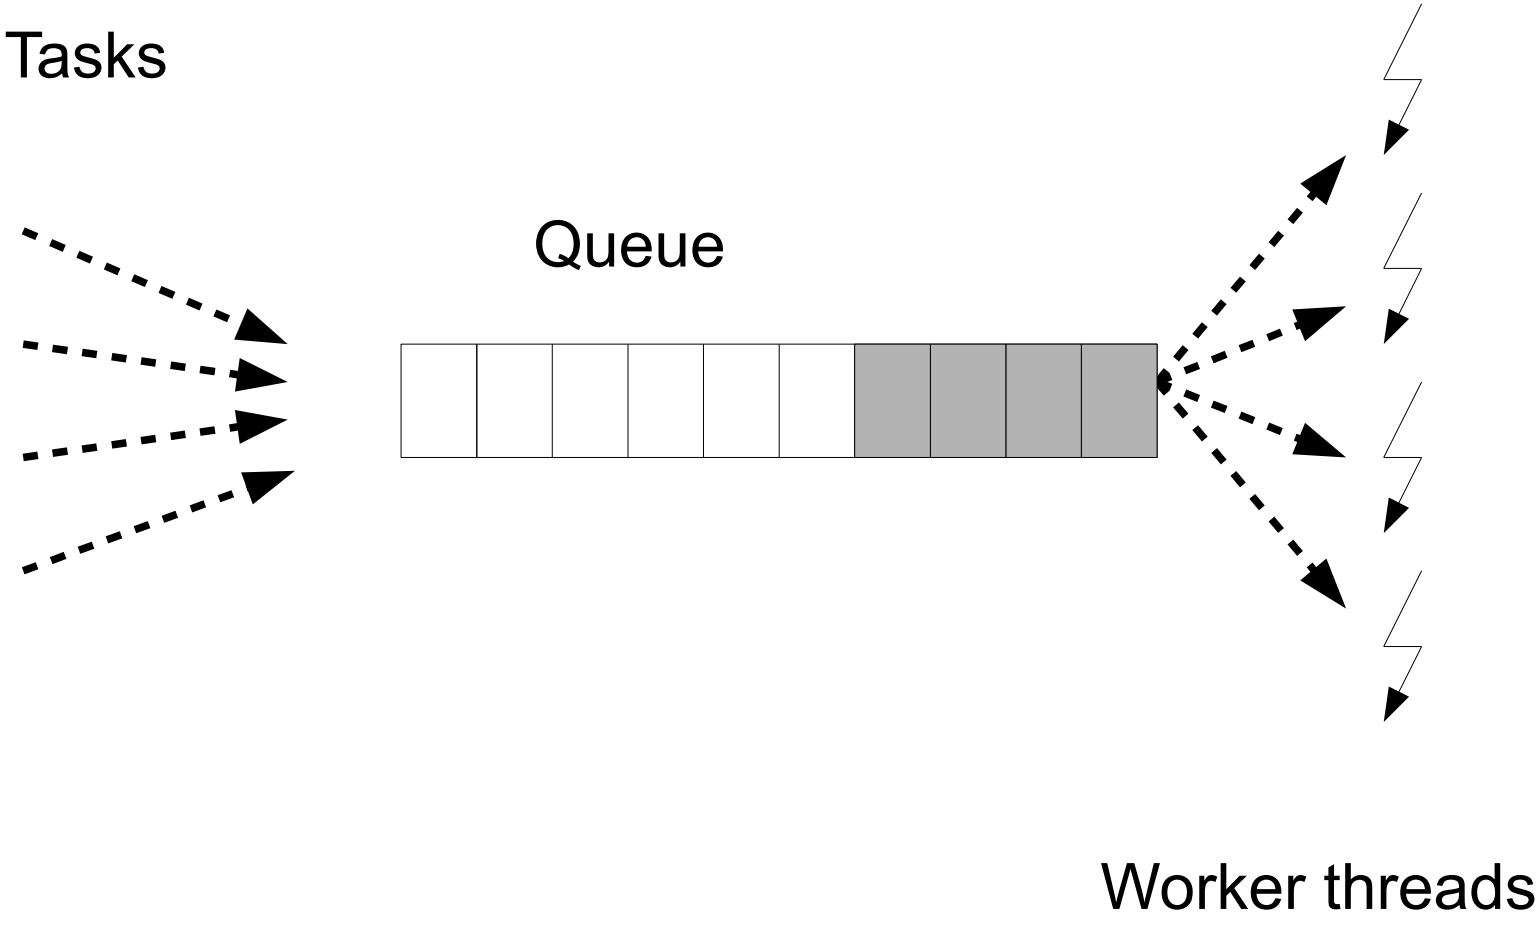

Thread Pool

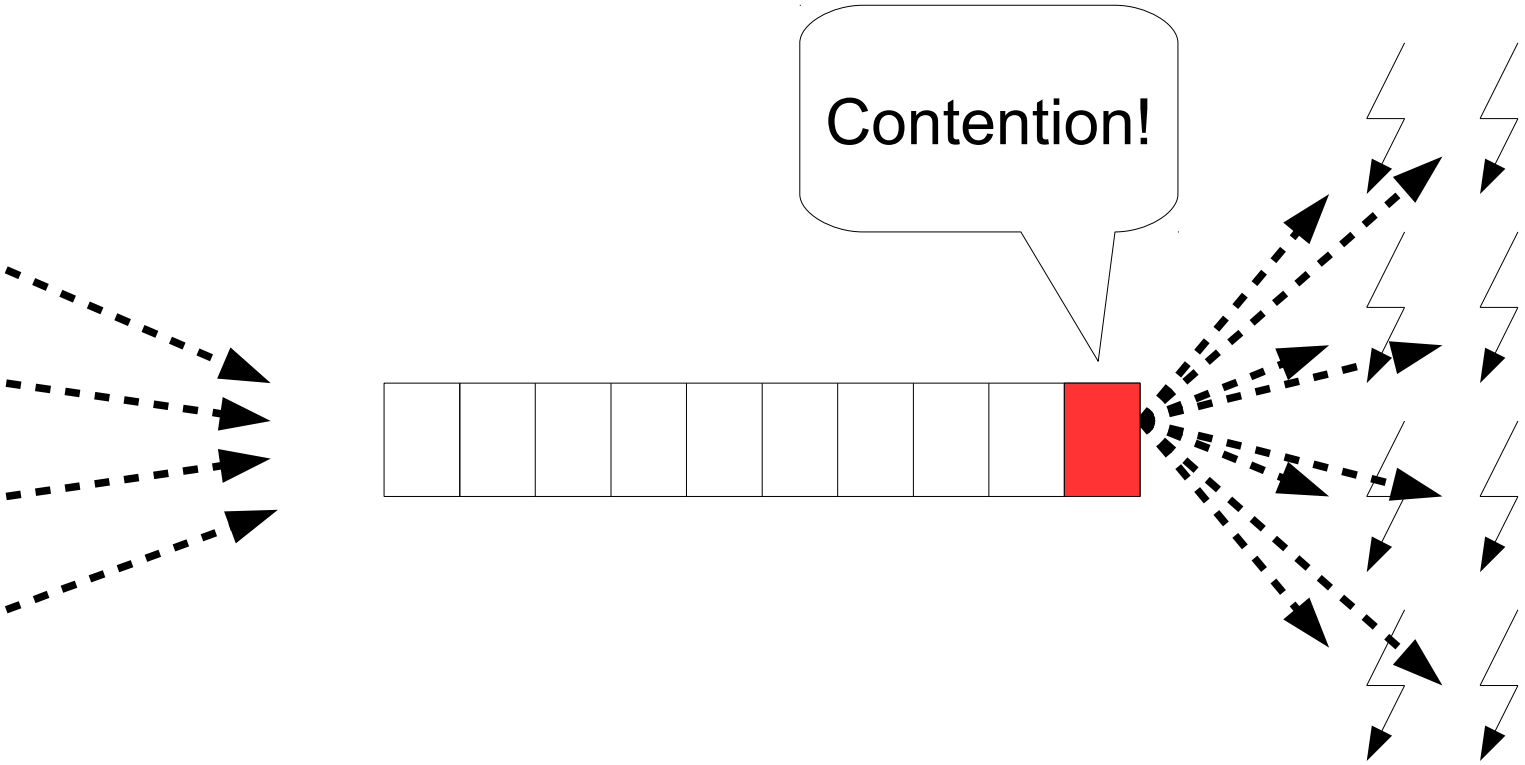

Fork/Join Thread Pool

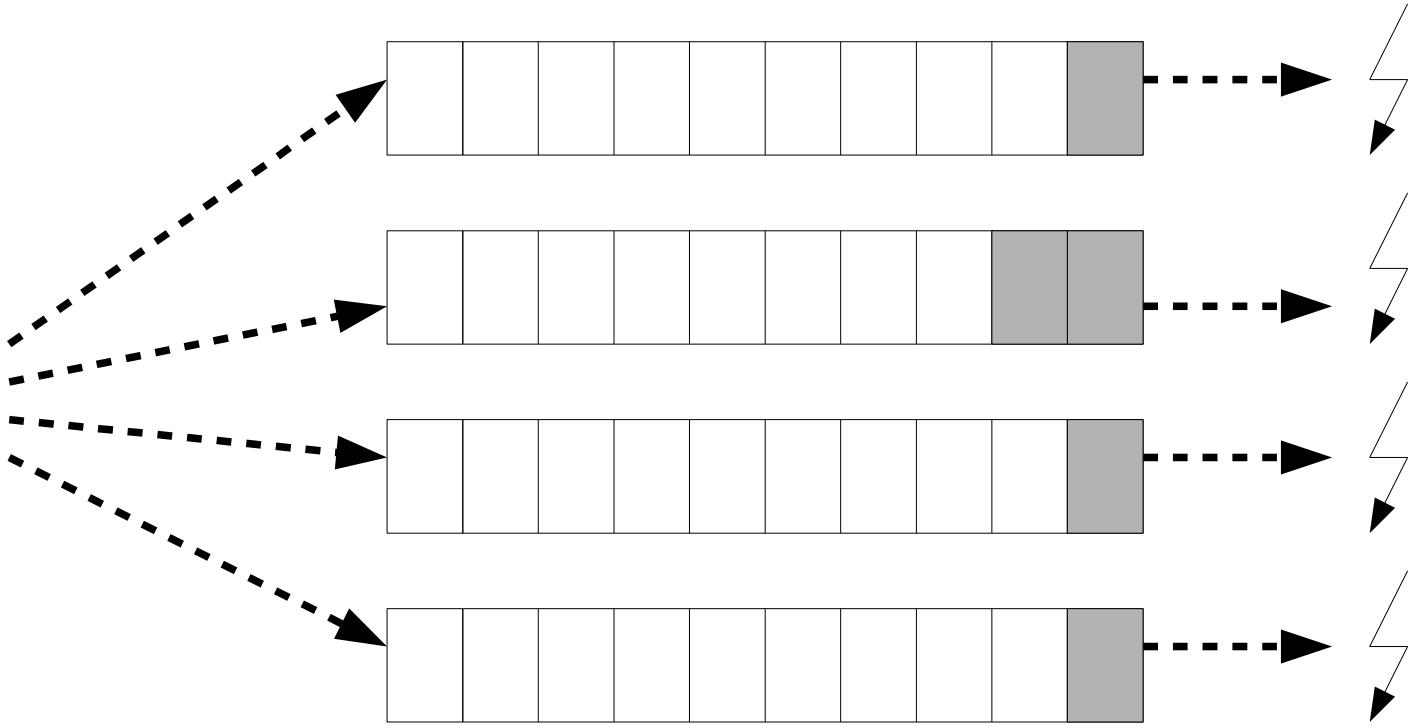

Fork/Join Thread Pool

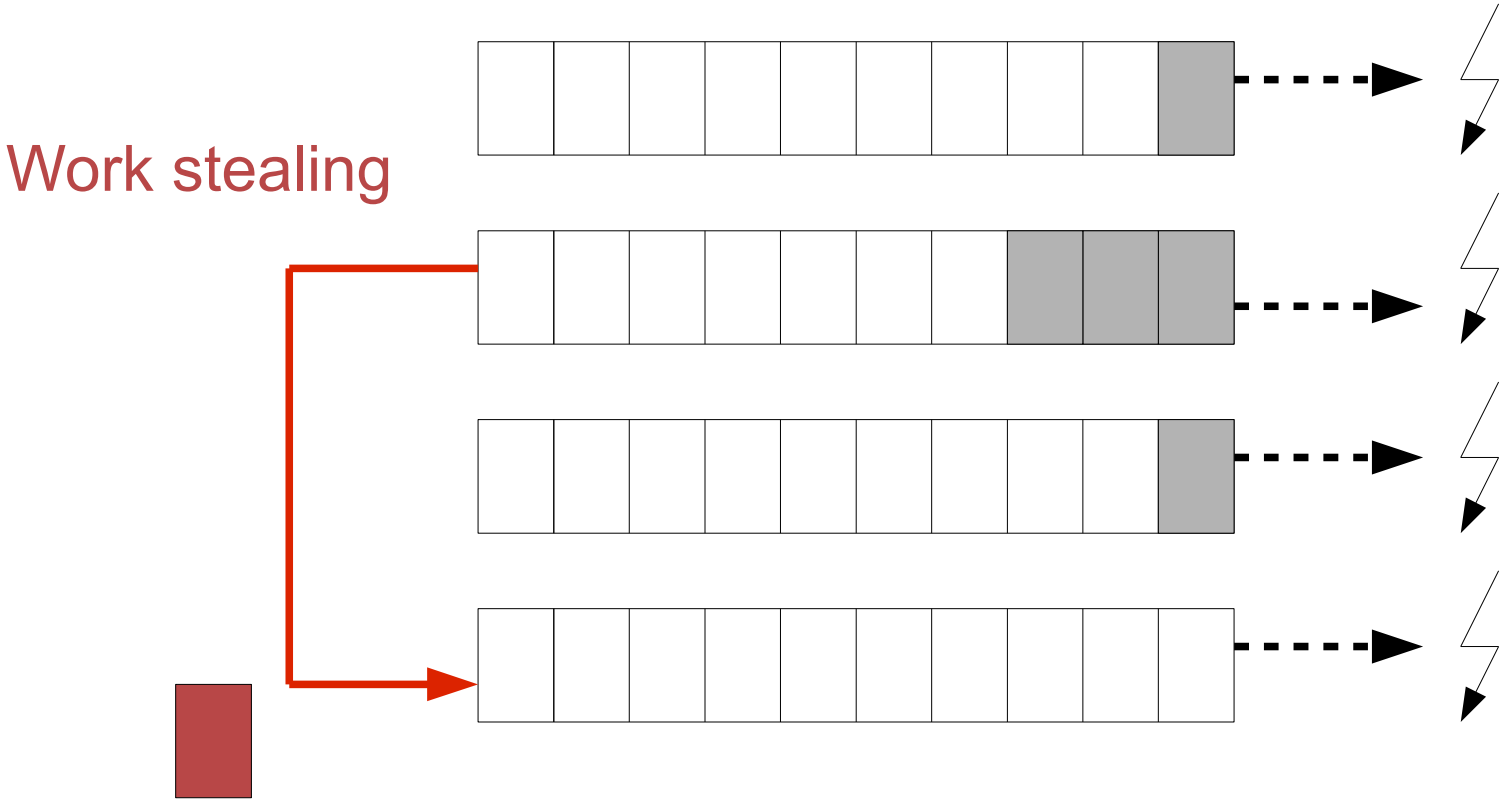

Async the Groovy way

```
task {
    calculation.process()
}
```

Async the Groovy way

```
def group = new NonDaemonPGroup(10)

group.task {
    calculation.process()
}
```

Async the Groovy way

```
group.task {->…}

group.task new Runnable() {…}

group.task new Callable<V>() {...}
```

Independent tasks

```
def group = new NonDaemonPGroup(10)

submissions.each {form →
    group.task {
        form.process()
    }
}
```

Dependent tasks

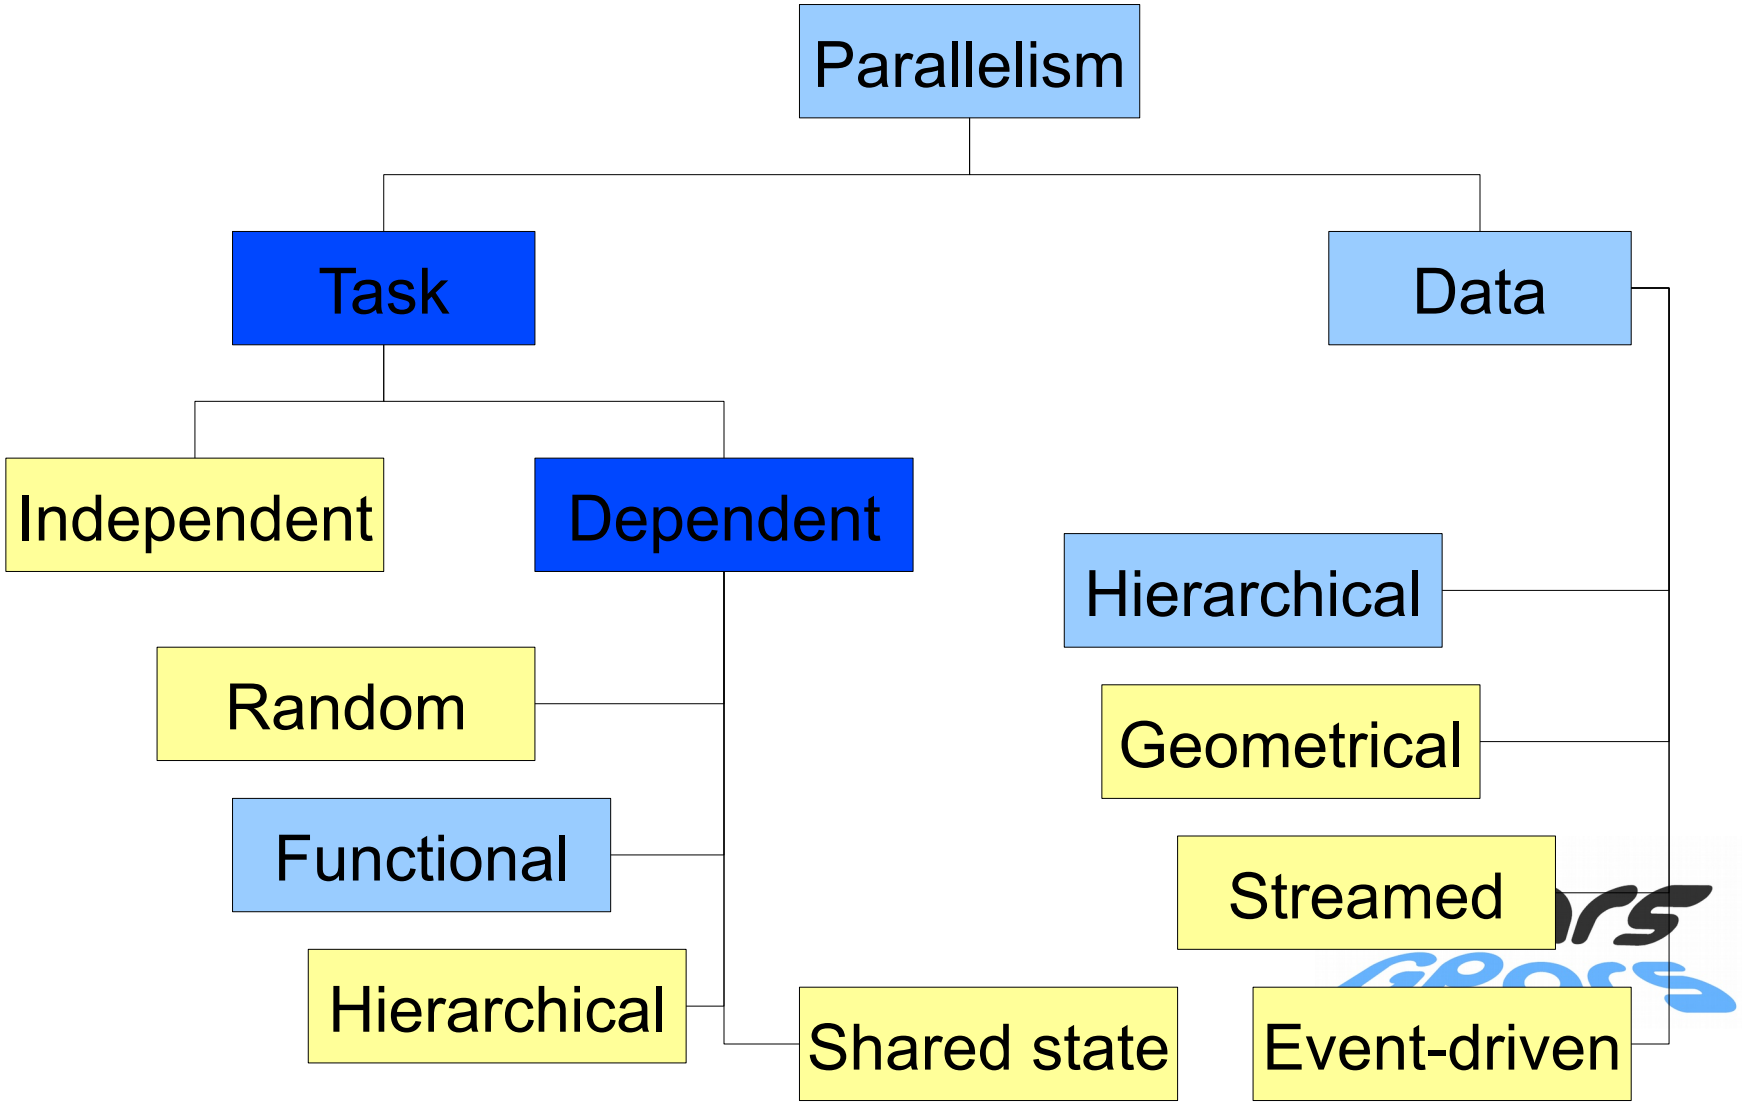

Dependent tasks

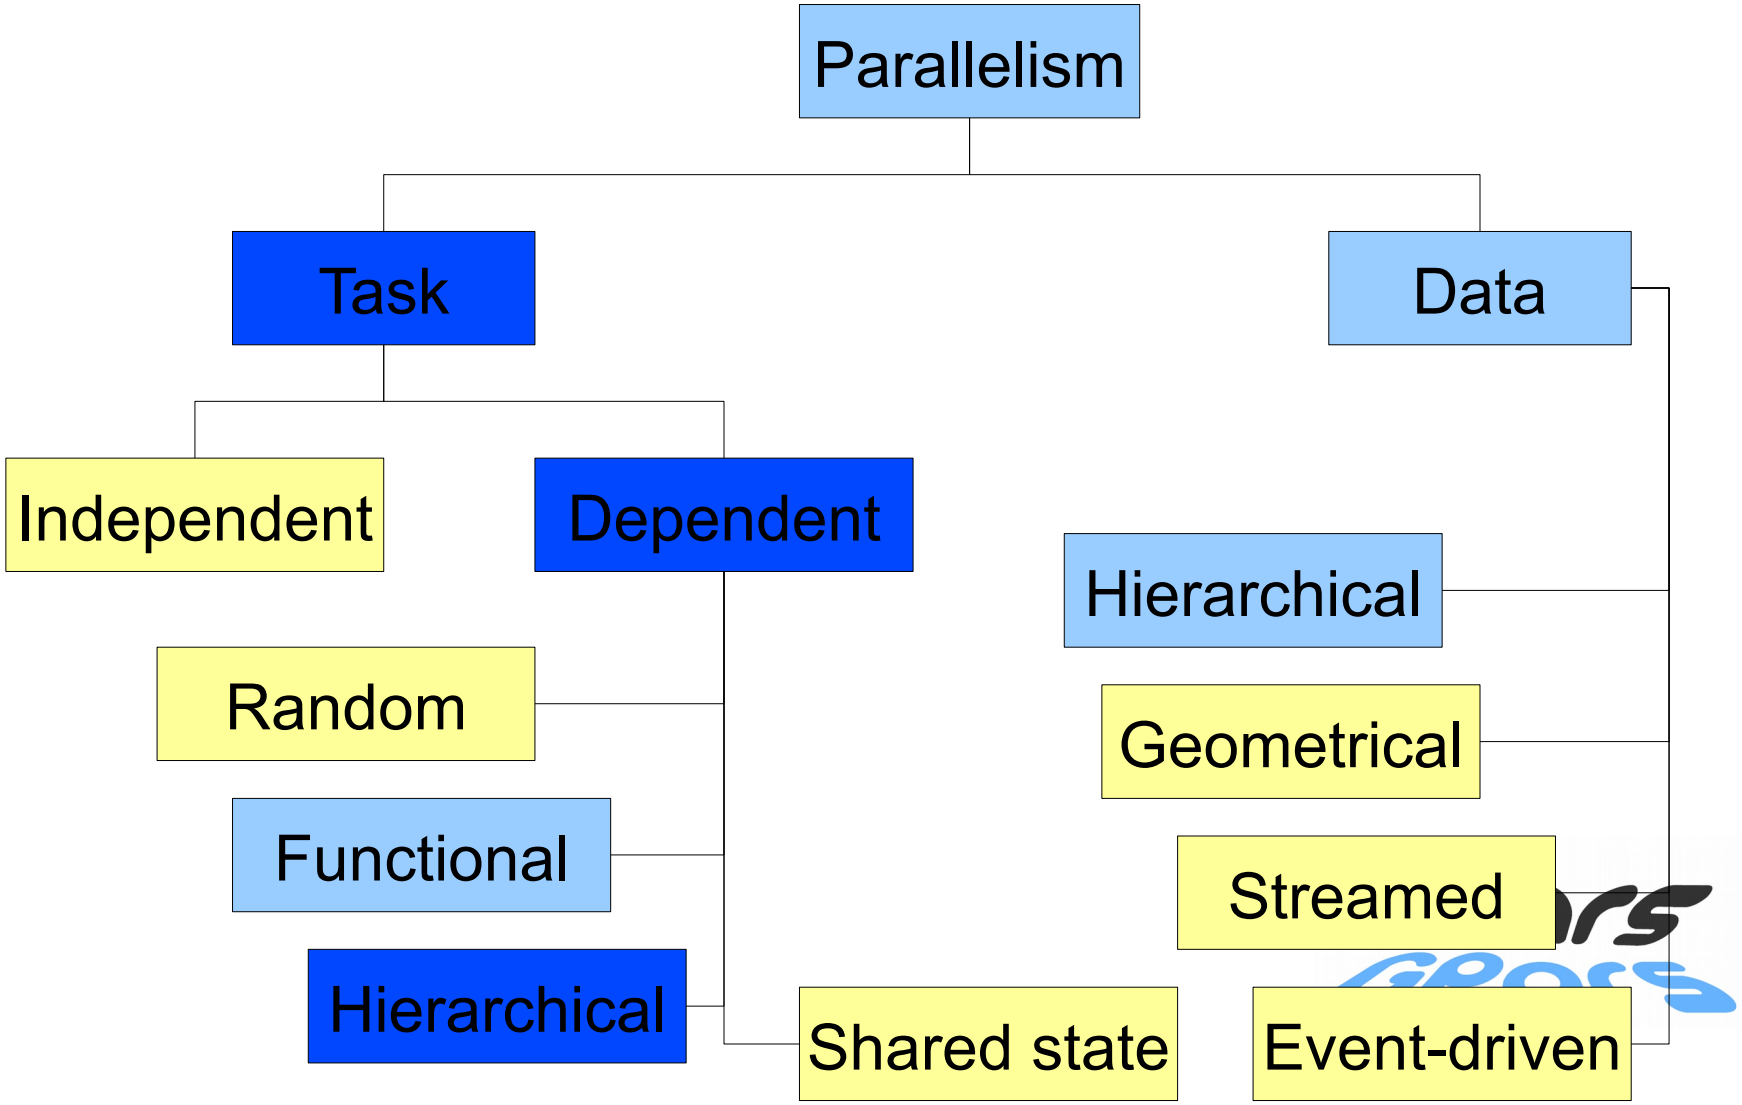

Hierarchical decomposition

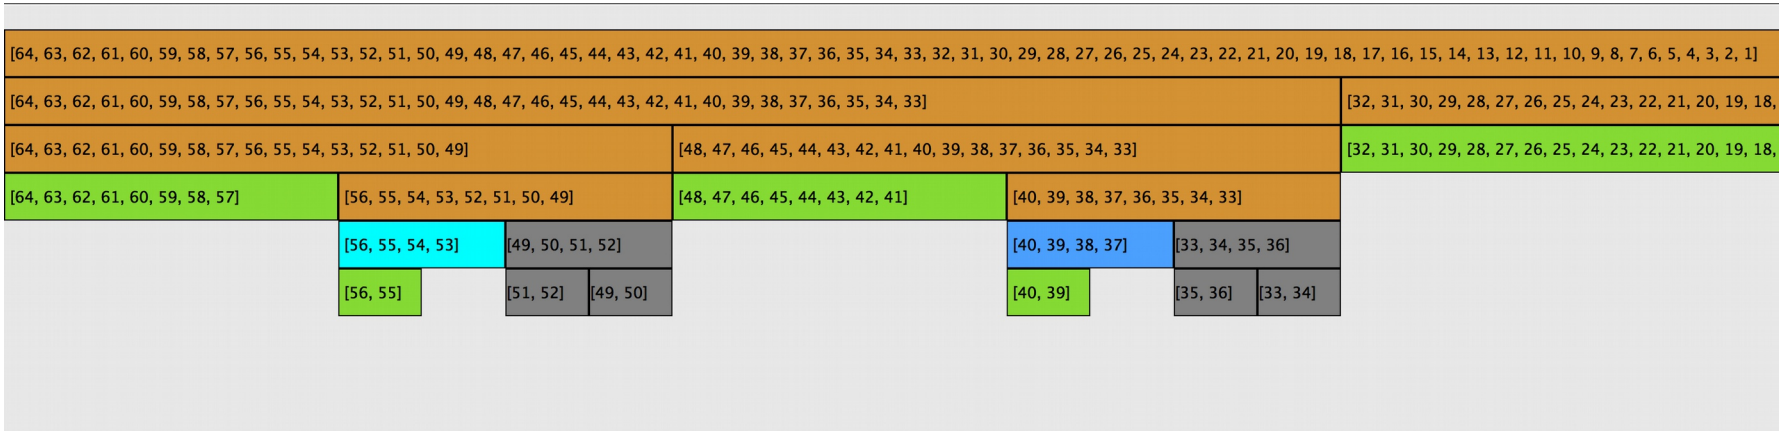

Fork/Join

- Solve hierarchical problems
  - Divide and conquer
  - Merge sort, Quick sort
  - Tree traversal
  - File scan / search
  - …

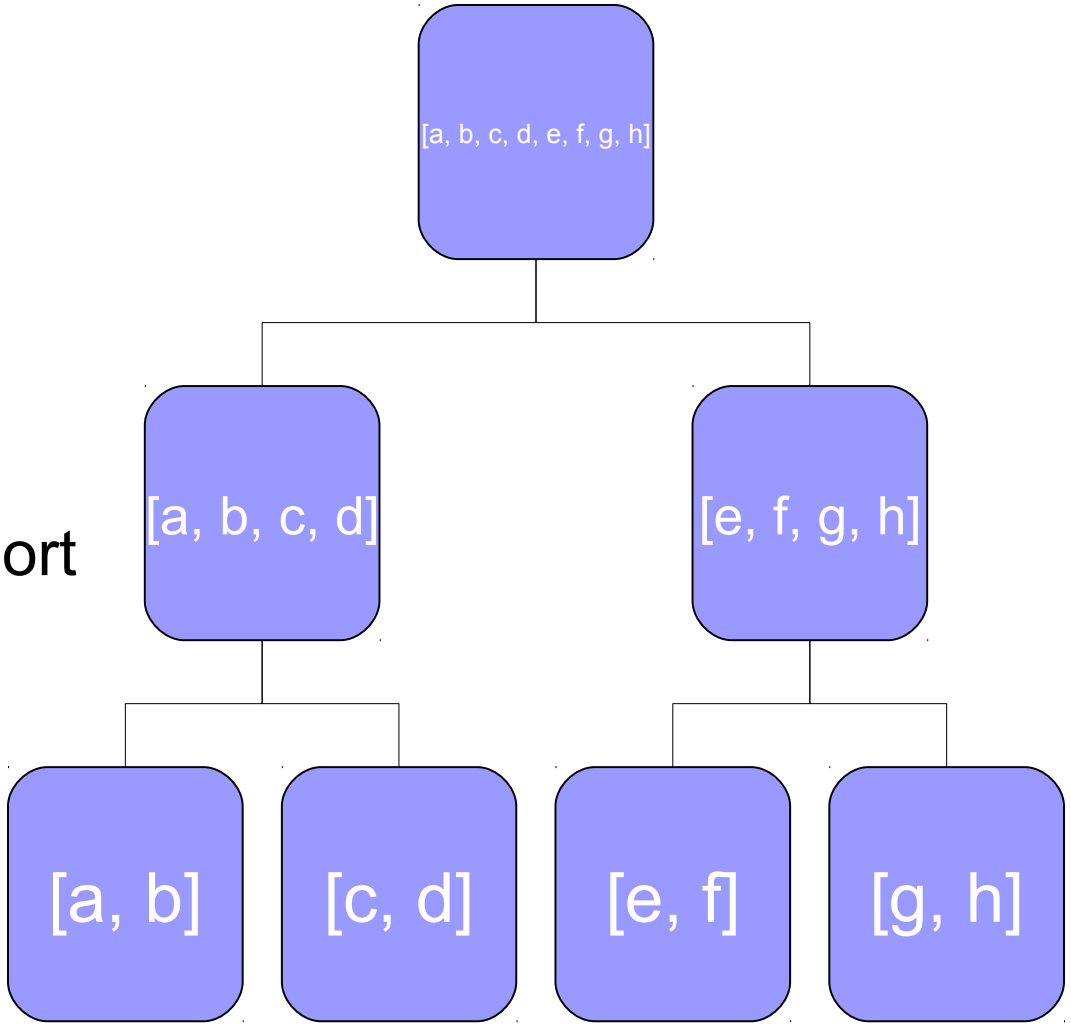

Fork/Join (GPars)

```
{currentDir ->
    long count = 0;
    currentDir.eachFile {
        if (it.isDirectory()) {
            forkOffChild it
        } else {
            count++
        }
    }
    return count + childrenResults.sum(0)
}
```

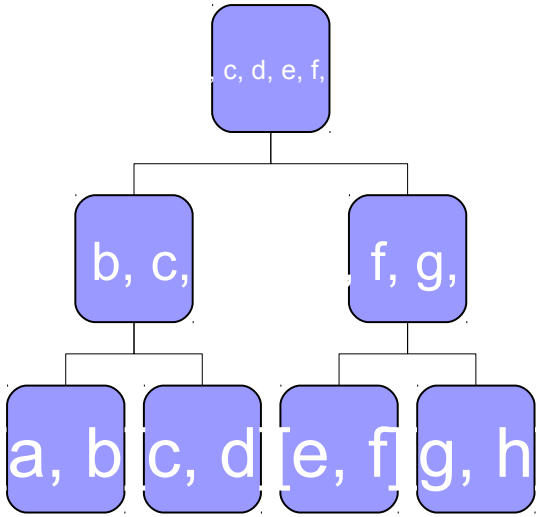

Waits for children without blocking the thread!

State sharing

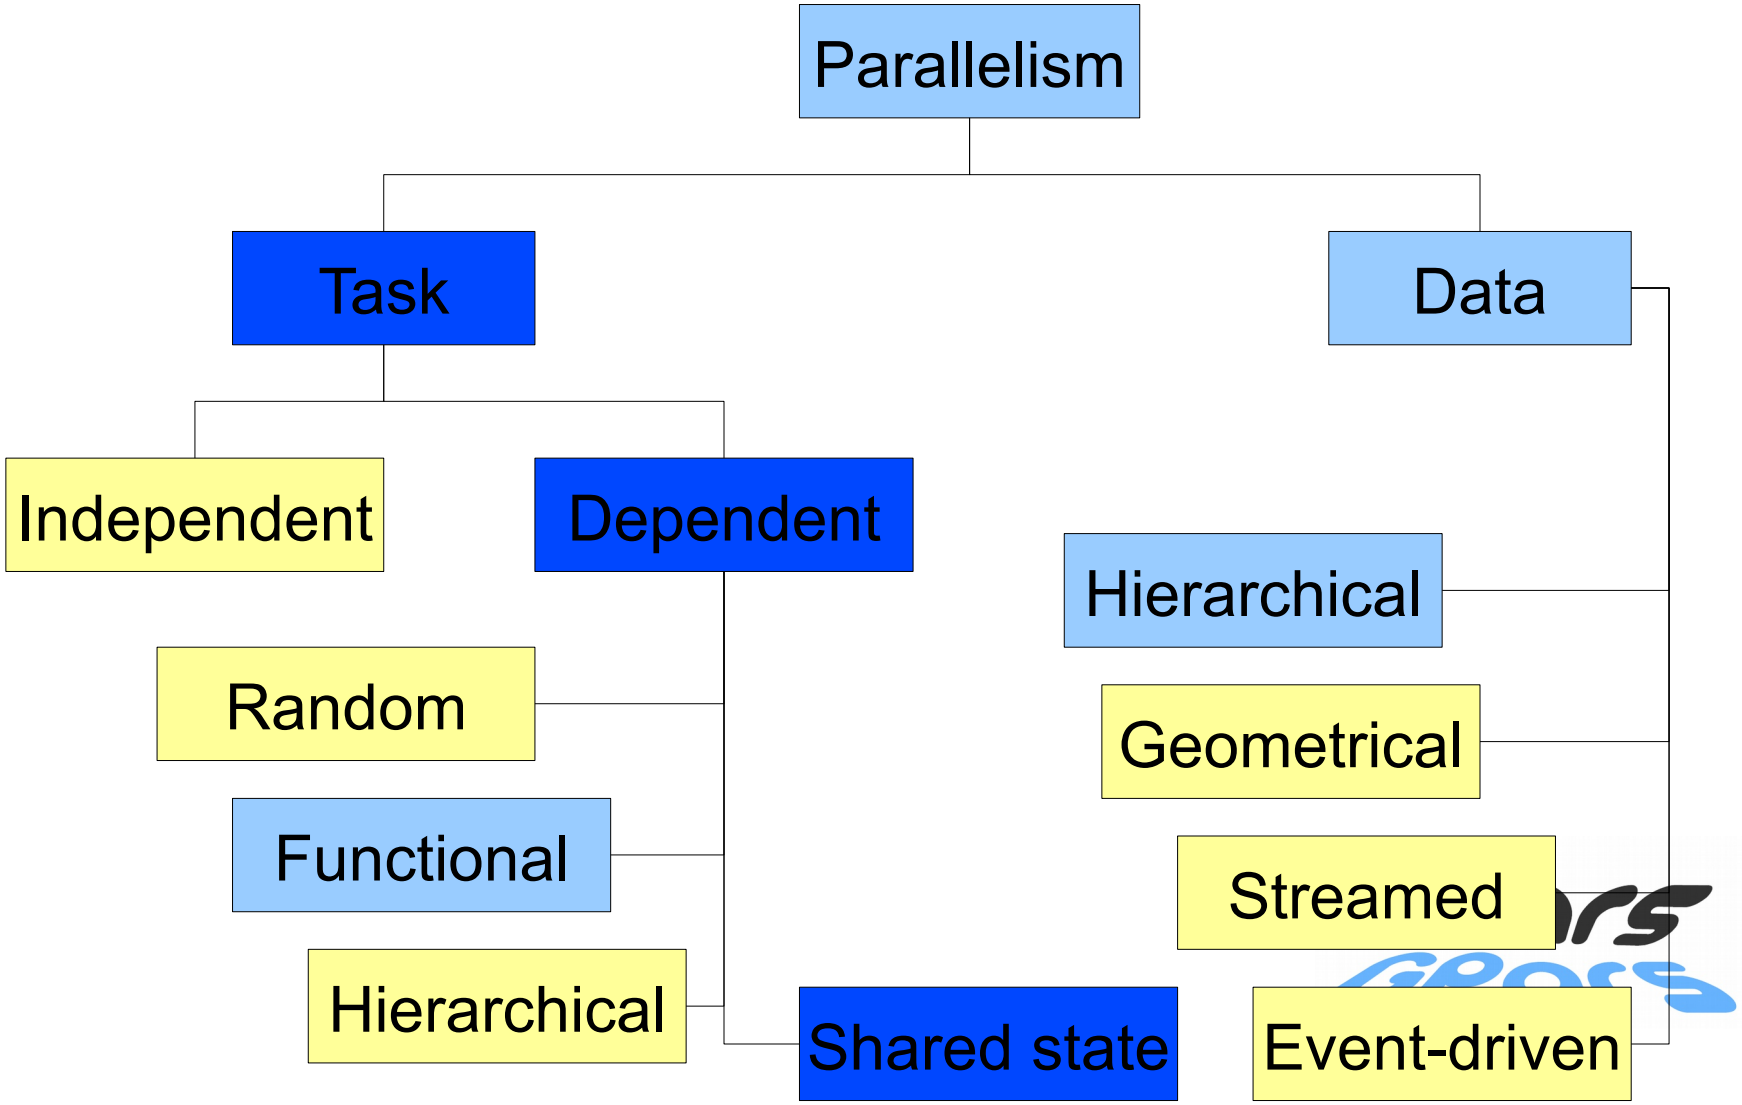

State sharing

```
List registrations = []
submissions.each {form →
    group.task {
        if (form.process().valid) {
            registrations << form
        }
    }
}
```

State sharing

Needs protection

```
List registrations = []
submissions.each {form →
    group.task {
        if (form.process().valid) {
            registrations << form
        }
    }
}
```

GPars

Shared Mutable State

Frequently misused

When really needed, use

- Agents
- Software Transactional Memory
- Locks

Agent

- Lock Shared Mutable State in a Safe

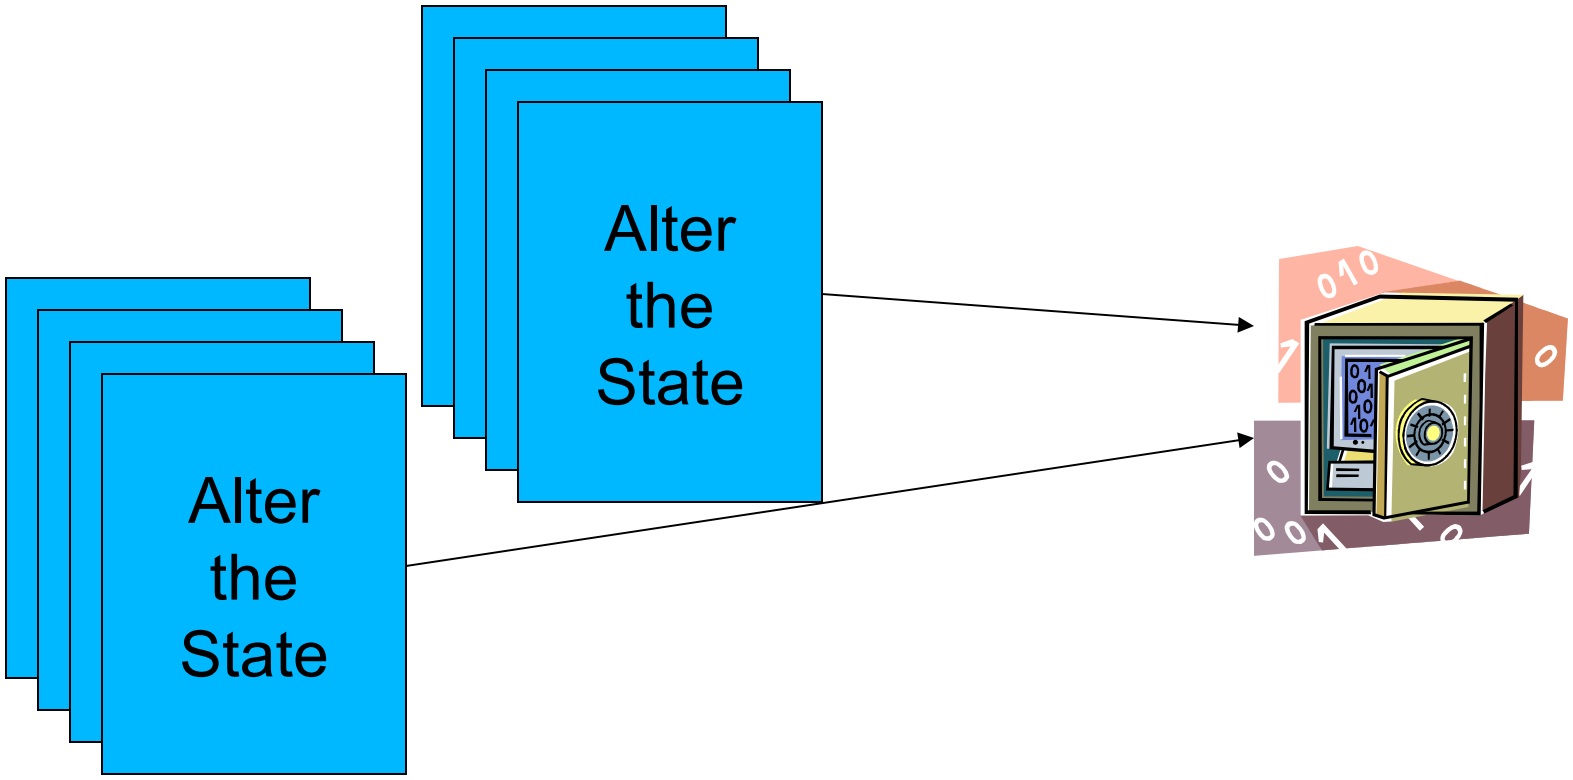

Agent inside

Double

Add 25

Increment

Message Queue

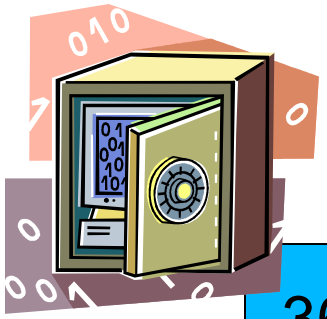

36

thread

Sharing through agents

```
Agent registrations = group.agent( [] )
submissions.each {form →
    group.task {
        if (form.process().valid) {
            registrations.send {it << form}
        }
    }
}
```

STM (Akka - Scala)

```
atomic {
    .. // do something within a transaction
}

atomic(maxNrOfRetries) { .. }
atomicReadOnly { .. }

atomically {
    .. // try to do something
} orElse {
    .. // if tx clash; try do do something else
}
```

Sample taken from Akka documentation - http://doc.akkasource.org/stm

Persistent Data Structures

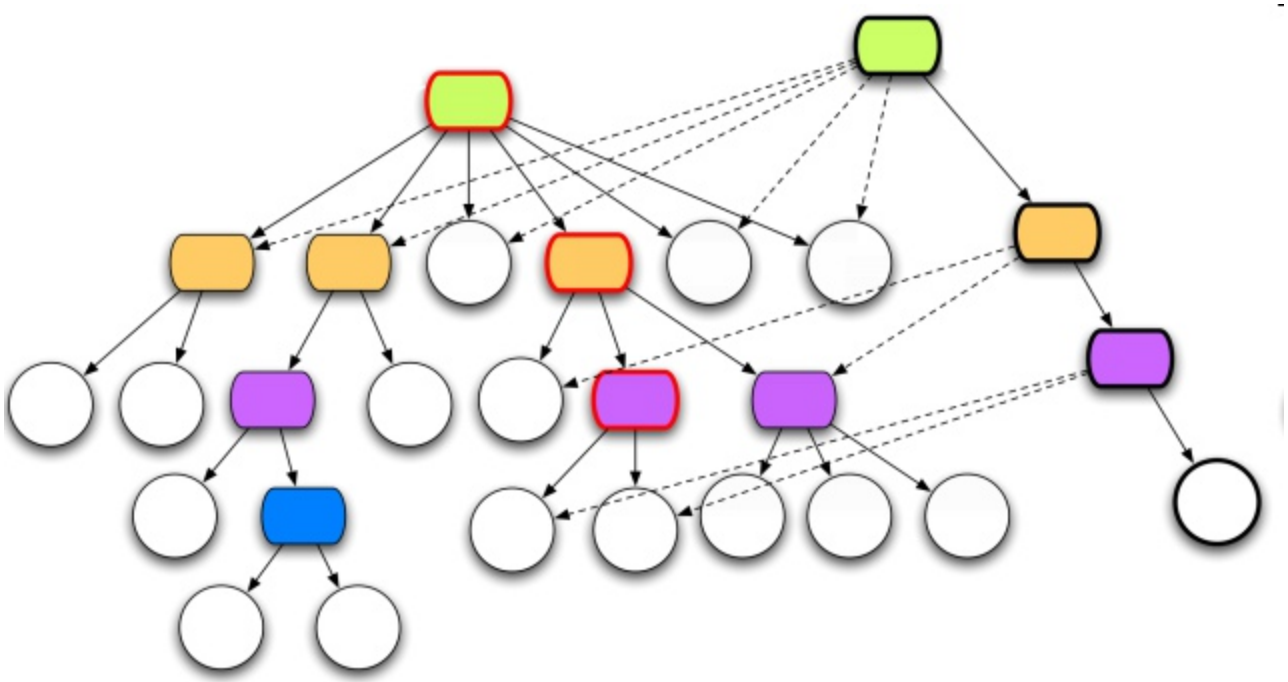

Illustration taken from Rich Hickey's presentation. Copyright Rich Hickey 2009

Random task dependency

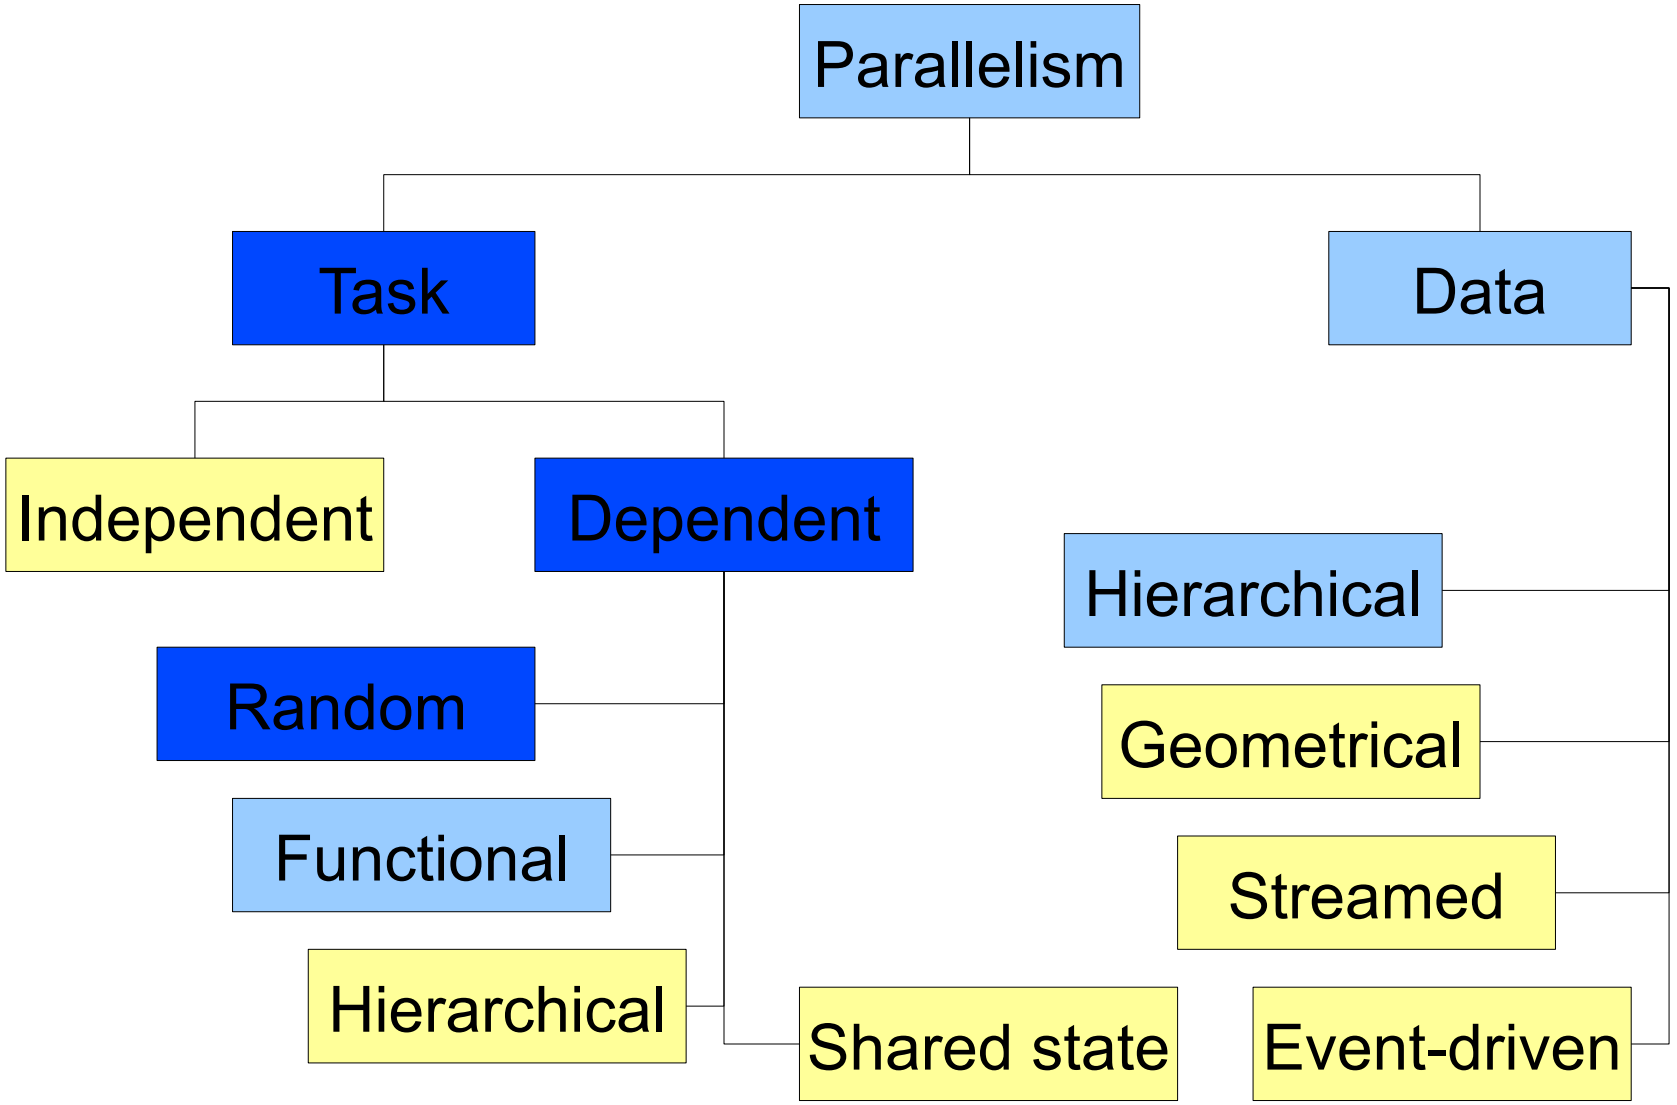

Dataflow Concurrency

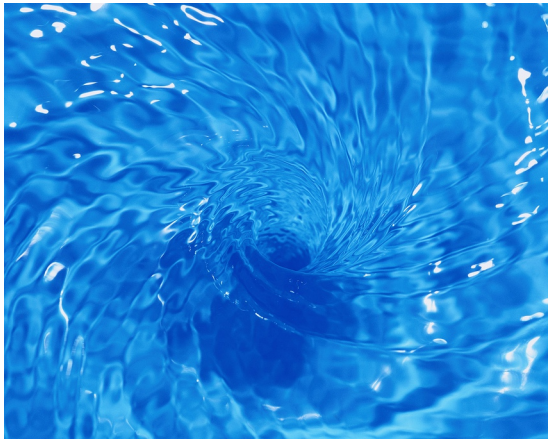

No race-conditions

No live-locks

Deterministic deadlocks

Dataflow Variables / Promises

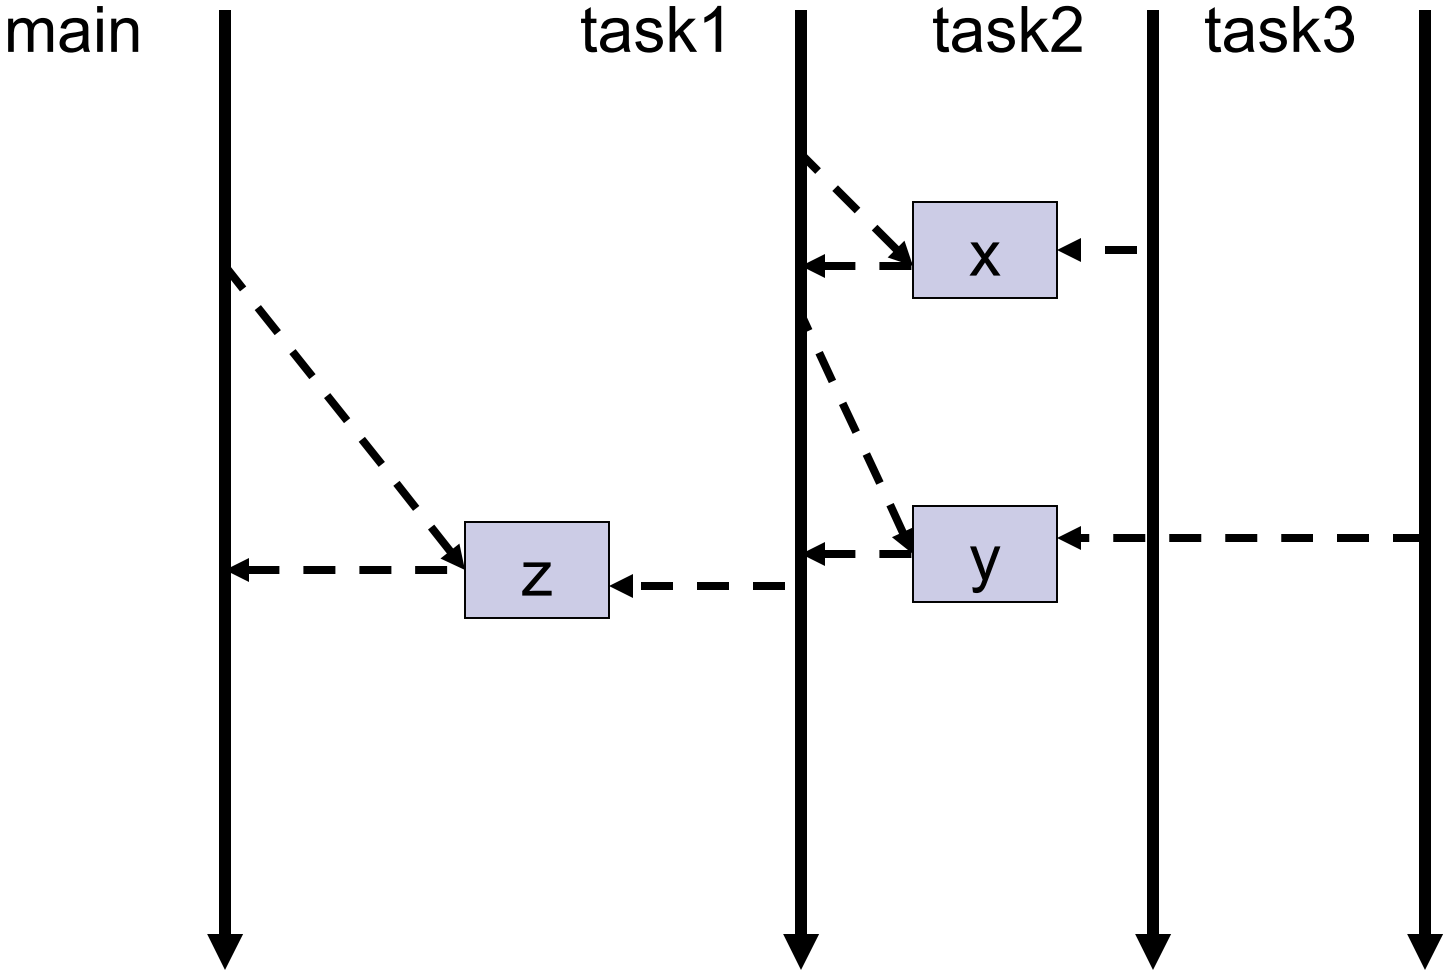

Dataflow Variables / Promises

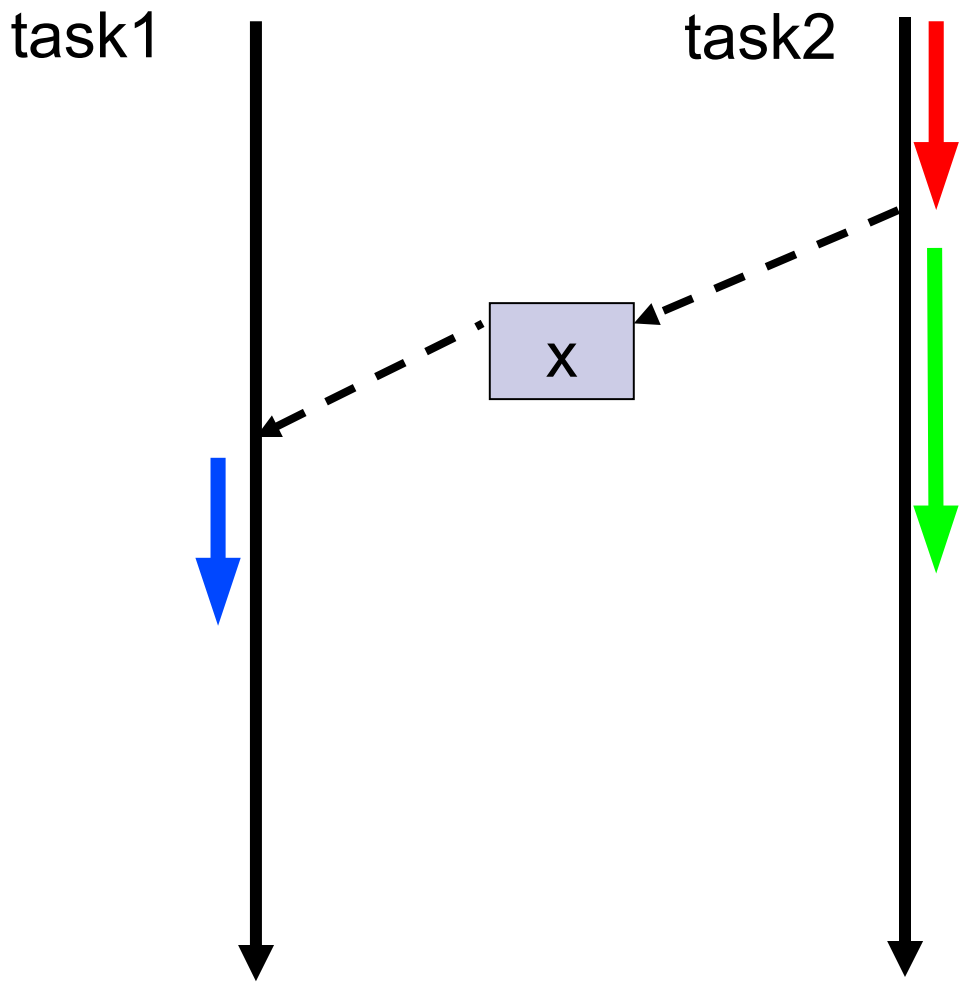

Dataflow Variables / Promises

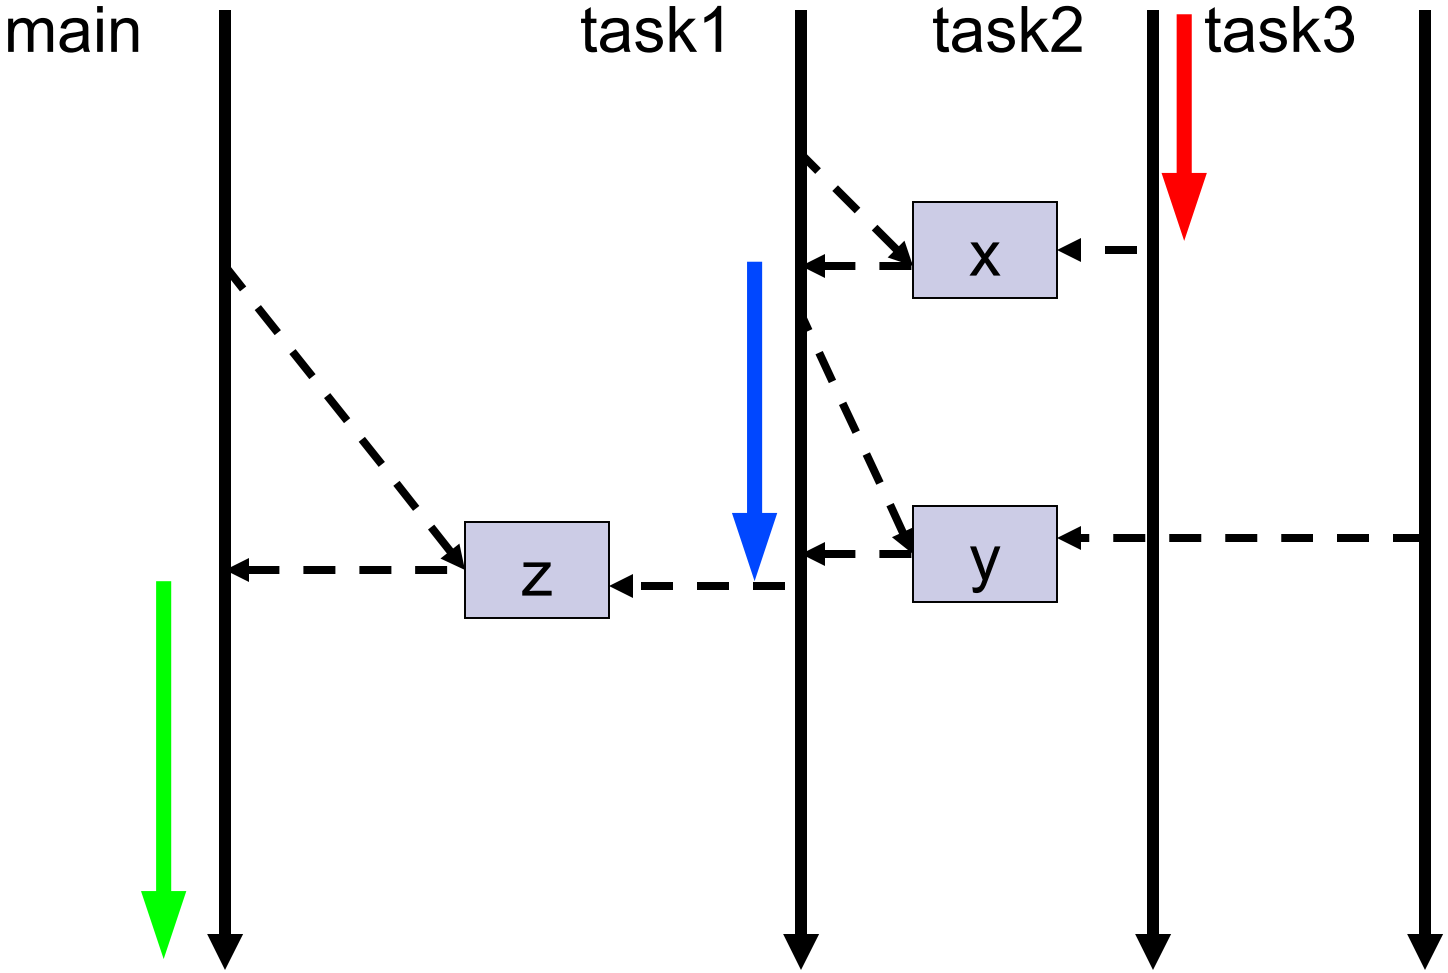

Synchronous Variables

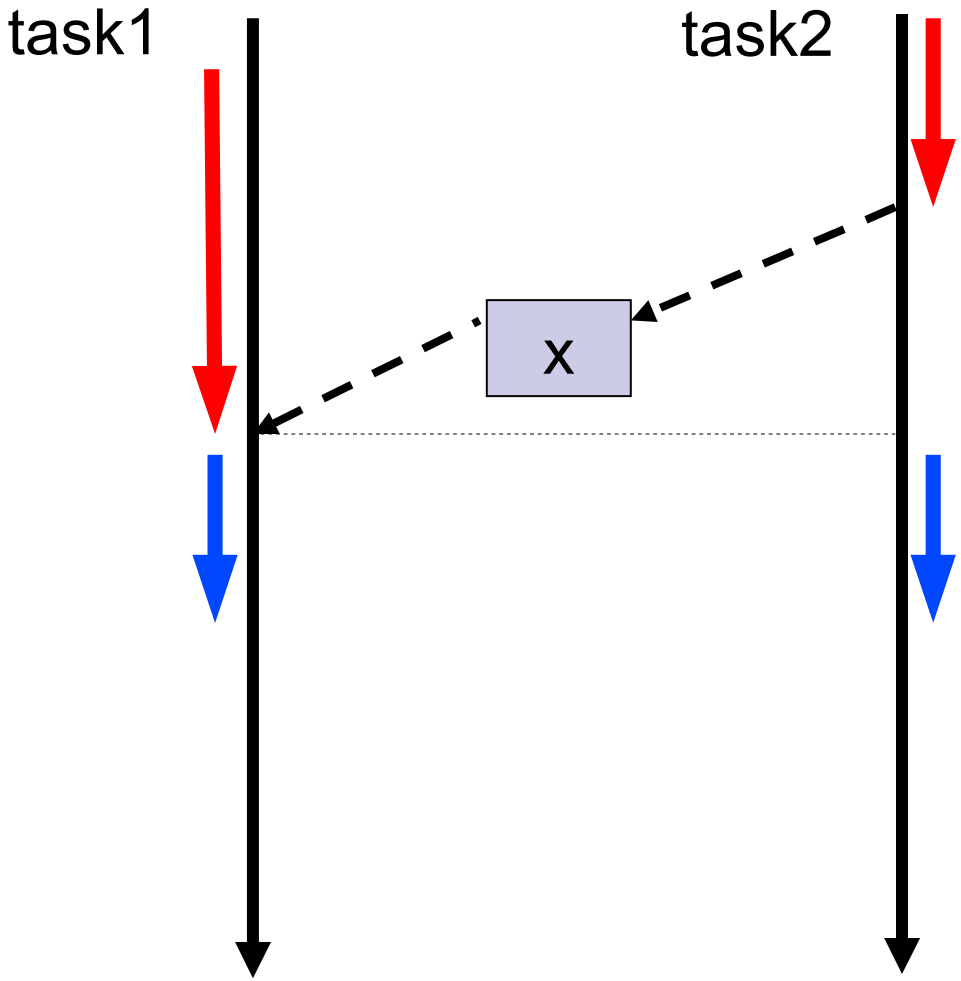

Synchronous Variables

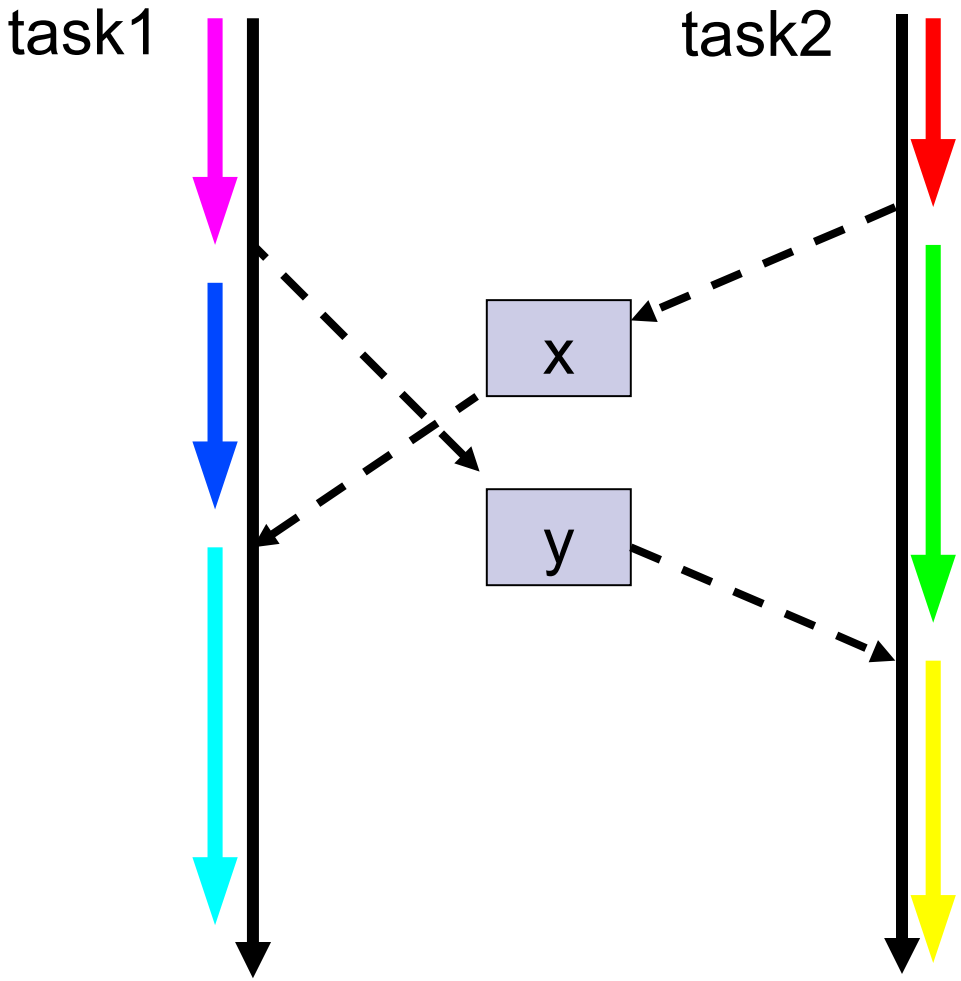

Synchronous Variables

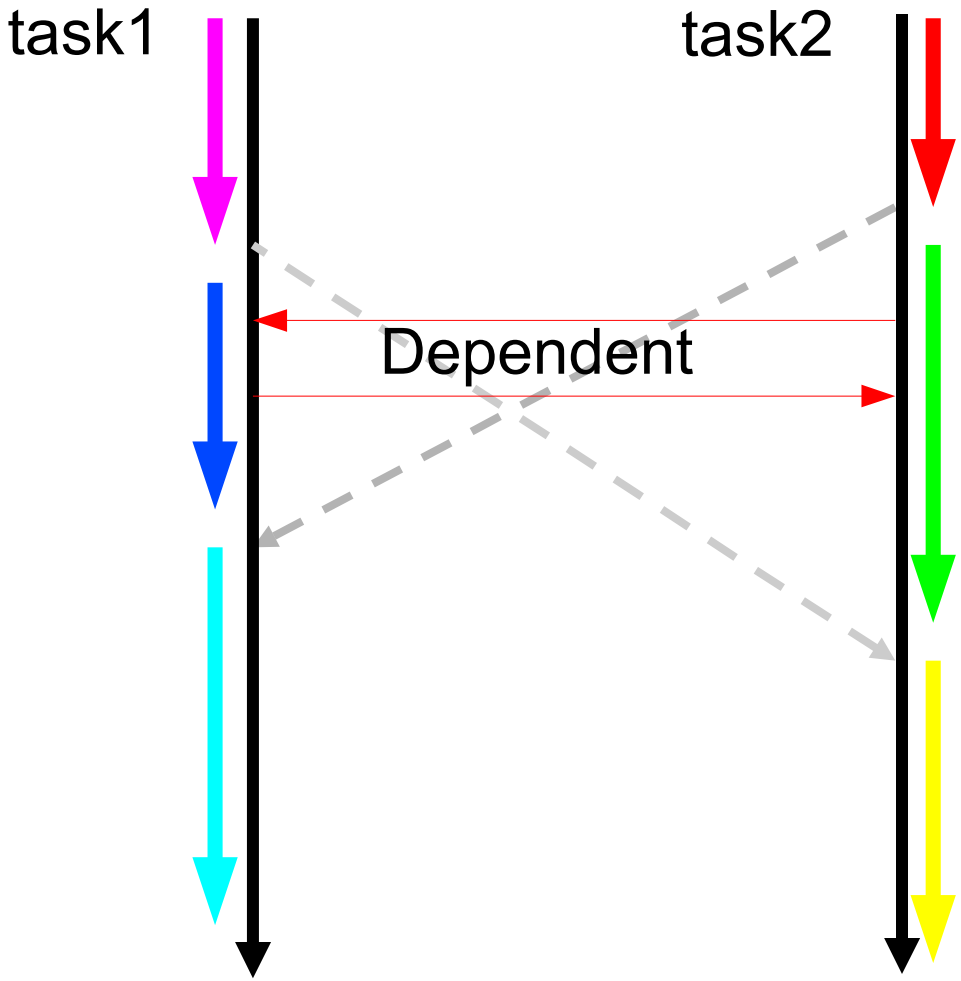

Promises to exchange data

```
task { z << x.val + y.val }
task { x << 10 }
task {
    println ”I am task 3”
    y << 5
}
assert 15 == z.val
```

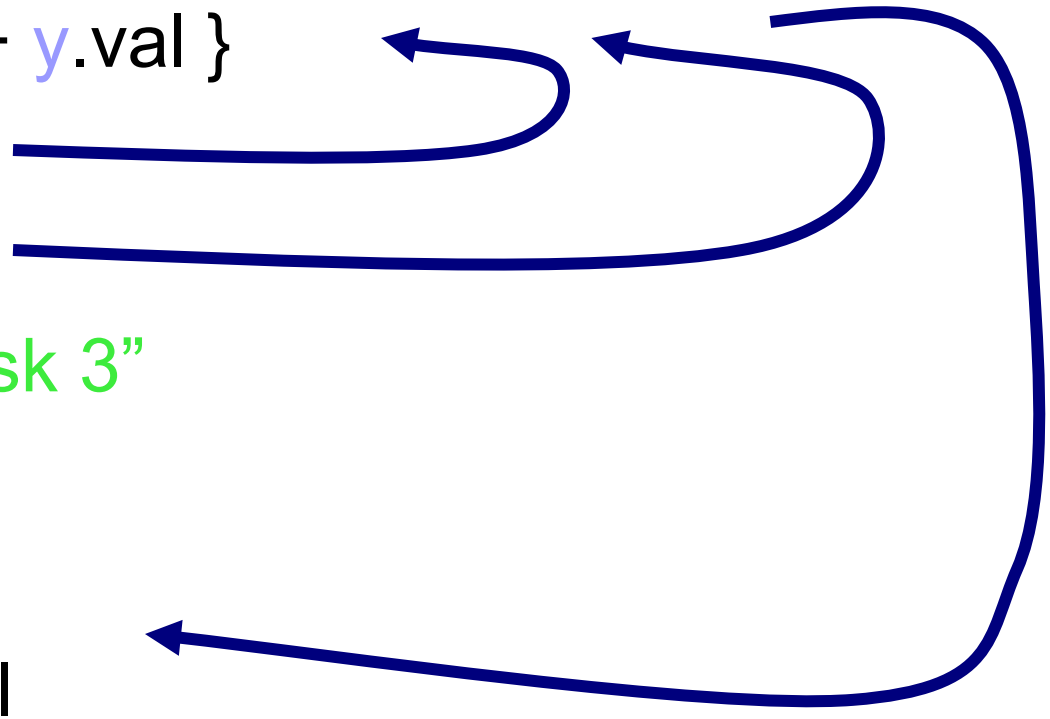

Glue tasks together

Promise c1 = task compile(module1)

Promise c2 = task compile(module2)

Glue tasks together

```
Promise c1 = task compile(module1)
Promise c2 = task compile(module2)

Promise j1 = c1.then {jar it}
Promise j2 = c2.then {jar it}
```

Glue tasks together

```
Promise c1 = task compile(module1)
Promise c2 = task compile(module2)

Promise j1 = c1.then {jar it}
Promise j2 = c2.then {jar it}
whenAllBound(j1, j2) {m1, m2 → deploy(m1, m2)}
j1.then {pushToRepo it}
```

Glue tasks together

```
Promise c1 = task compile(module1)
Promise c2 = task compile(module2)

Promise j1 = c1.then {jar it}
Promise j2 = c2.then {jar it}
whenAllBound(j1, j2) {m1, m2 → deploy(m1, m2)}
j1.then {pushToRepo it}
iWillSendEmailWhenJarred(j1)
```

Chaining promises

```
def h1 = download('url') then {text → text.trim()} then hash
```

Chaining promises

```
def h1 = download('url') then {text → text.trim()} then hash
```

```
def h1 = download('url') | {text → text.trim()} | hash
```

Error handling

```
url.then(download)
    .then(calculateHash)
    .then(formatResult)
    .then(printResult, printError)
    .then(sendNotificationEmail);
```

Lazy promises

Only calculated when needed the first time

```
def mostPopularLang = new LazyDataflowVariable({->
        return longLastingCalculation()
    }
})
```

Dataflow Variables / Promises

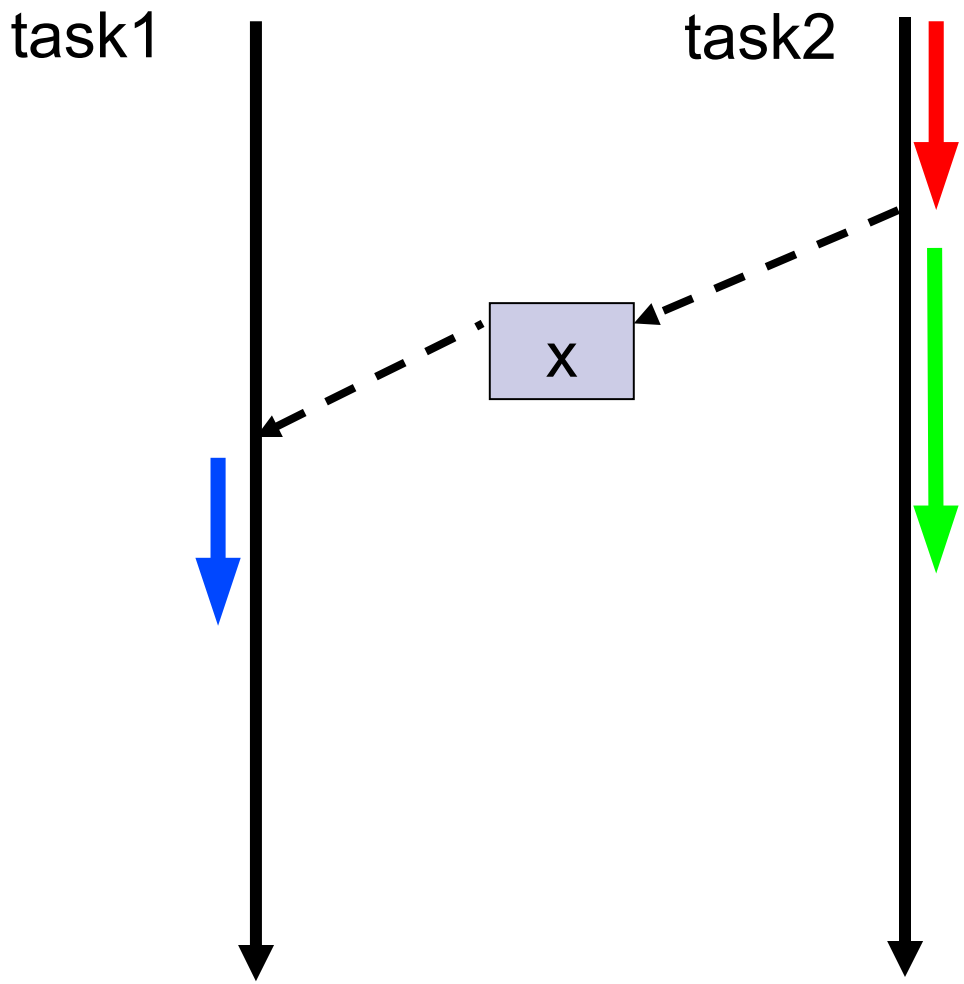

Dataflow Channels

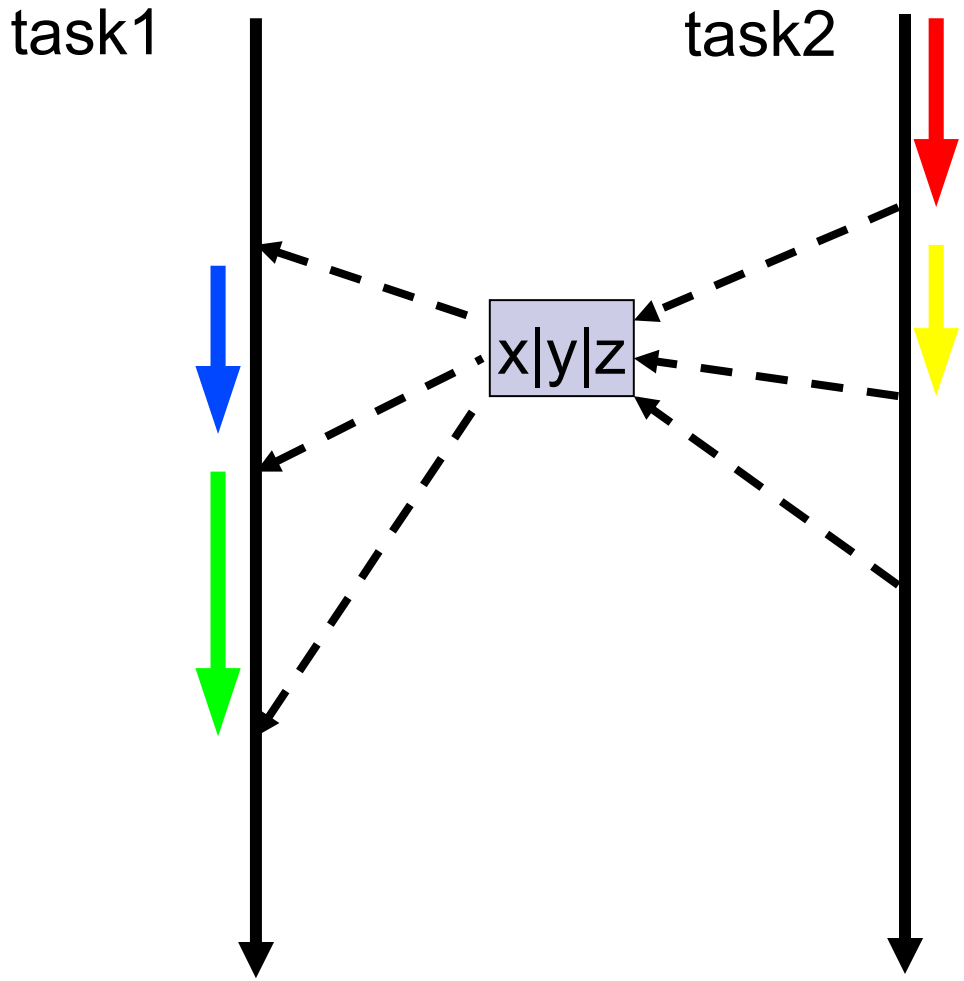

Synchronous Channels

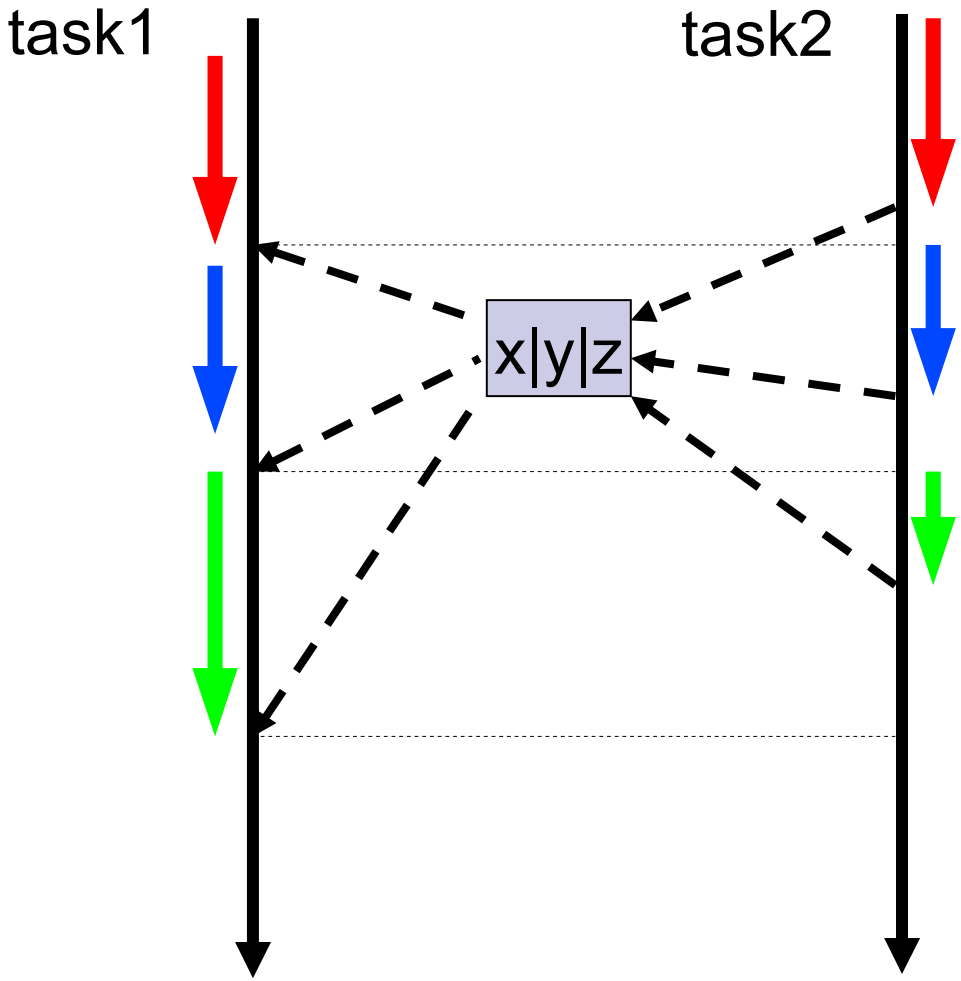

Tasks with progress indication

```
List<Promise> forms=submissions.collect {form →
    group.task {
        def result = form.process()
        progressQueue << 1
        if (result.valid) {
            return form
        }
    }
}
```

Channel Selection

```
Select alt = group.select(validForms, invalidForms)

SelectResult selectResult = alt.select()  //alt.prioritySelect()

switch (selectResult.index) {
        case 0: registrations << selectResult.value; break
        case 1: ...
}
```

Tasks as processes

```
group.task {
        doStuff()
        logChannel << 'initialized'
        def result = doWork(workQueue.val)
        if (result.isError) errors << result
        else results << result
        logChannel << 'finished'
}
```

A CSP flavour

Communicating Sequential Processes

Focus on composable processes more than on data

Processes

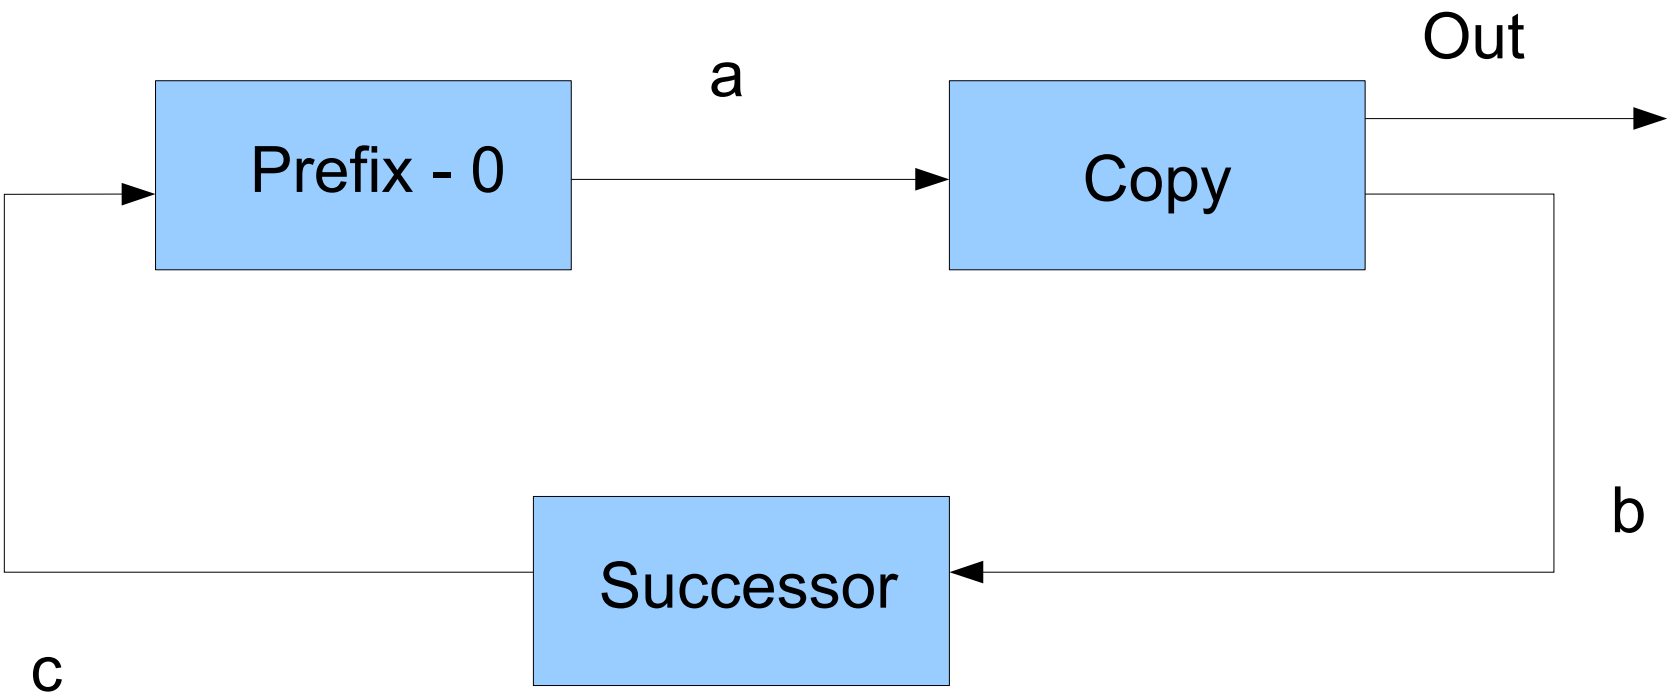

Processes

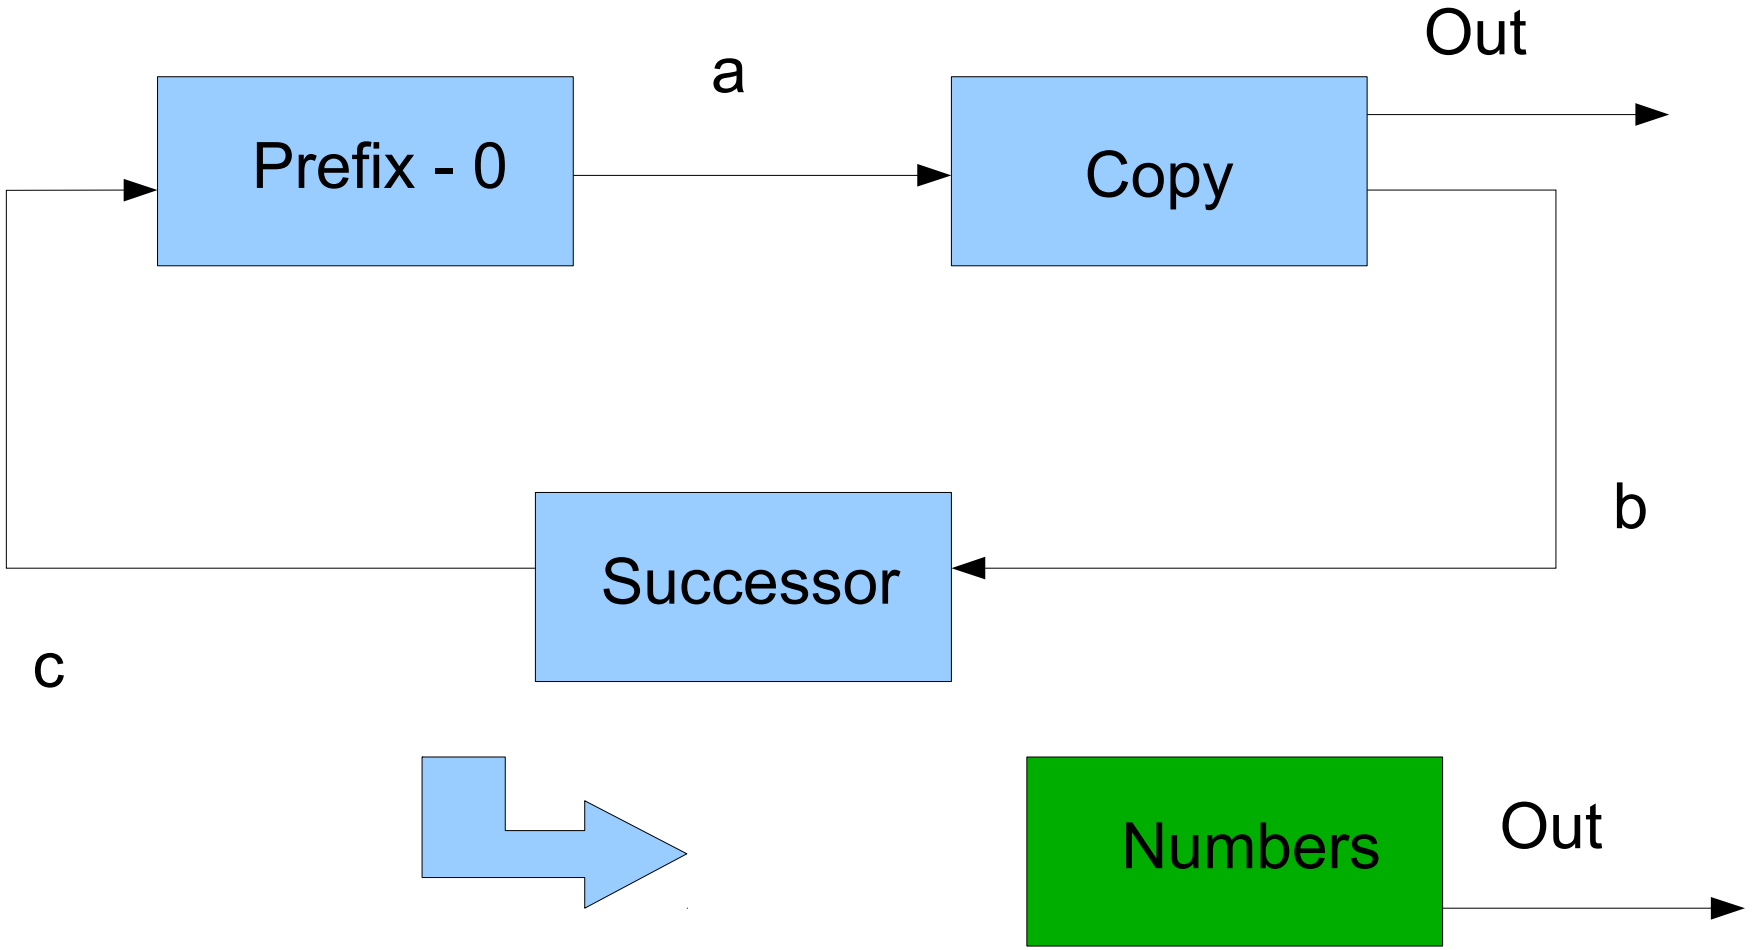

Processes

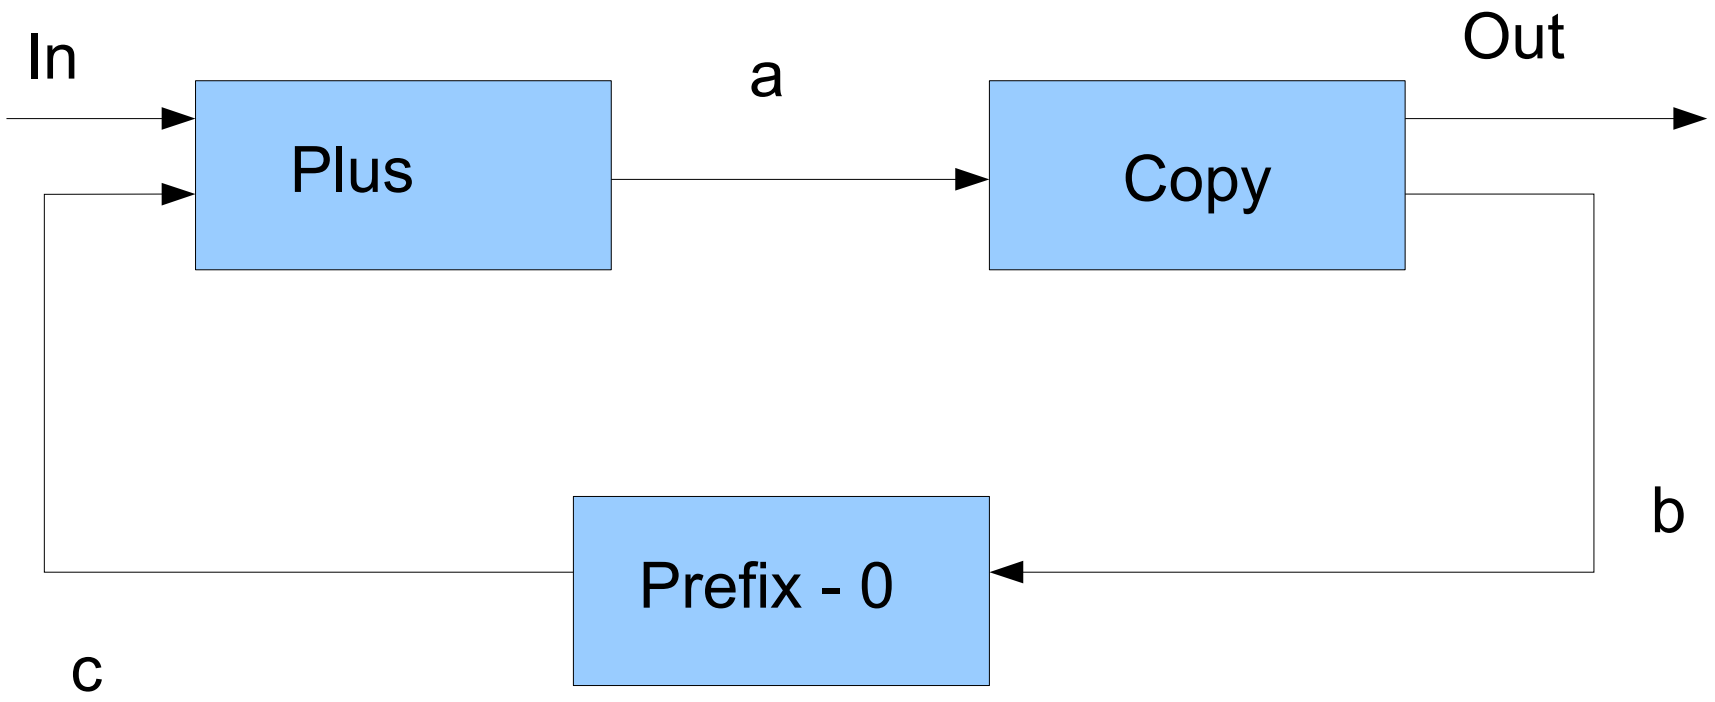

Processes

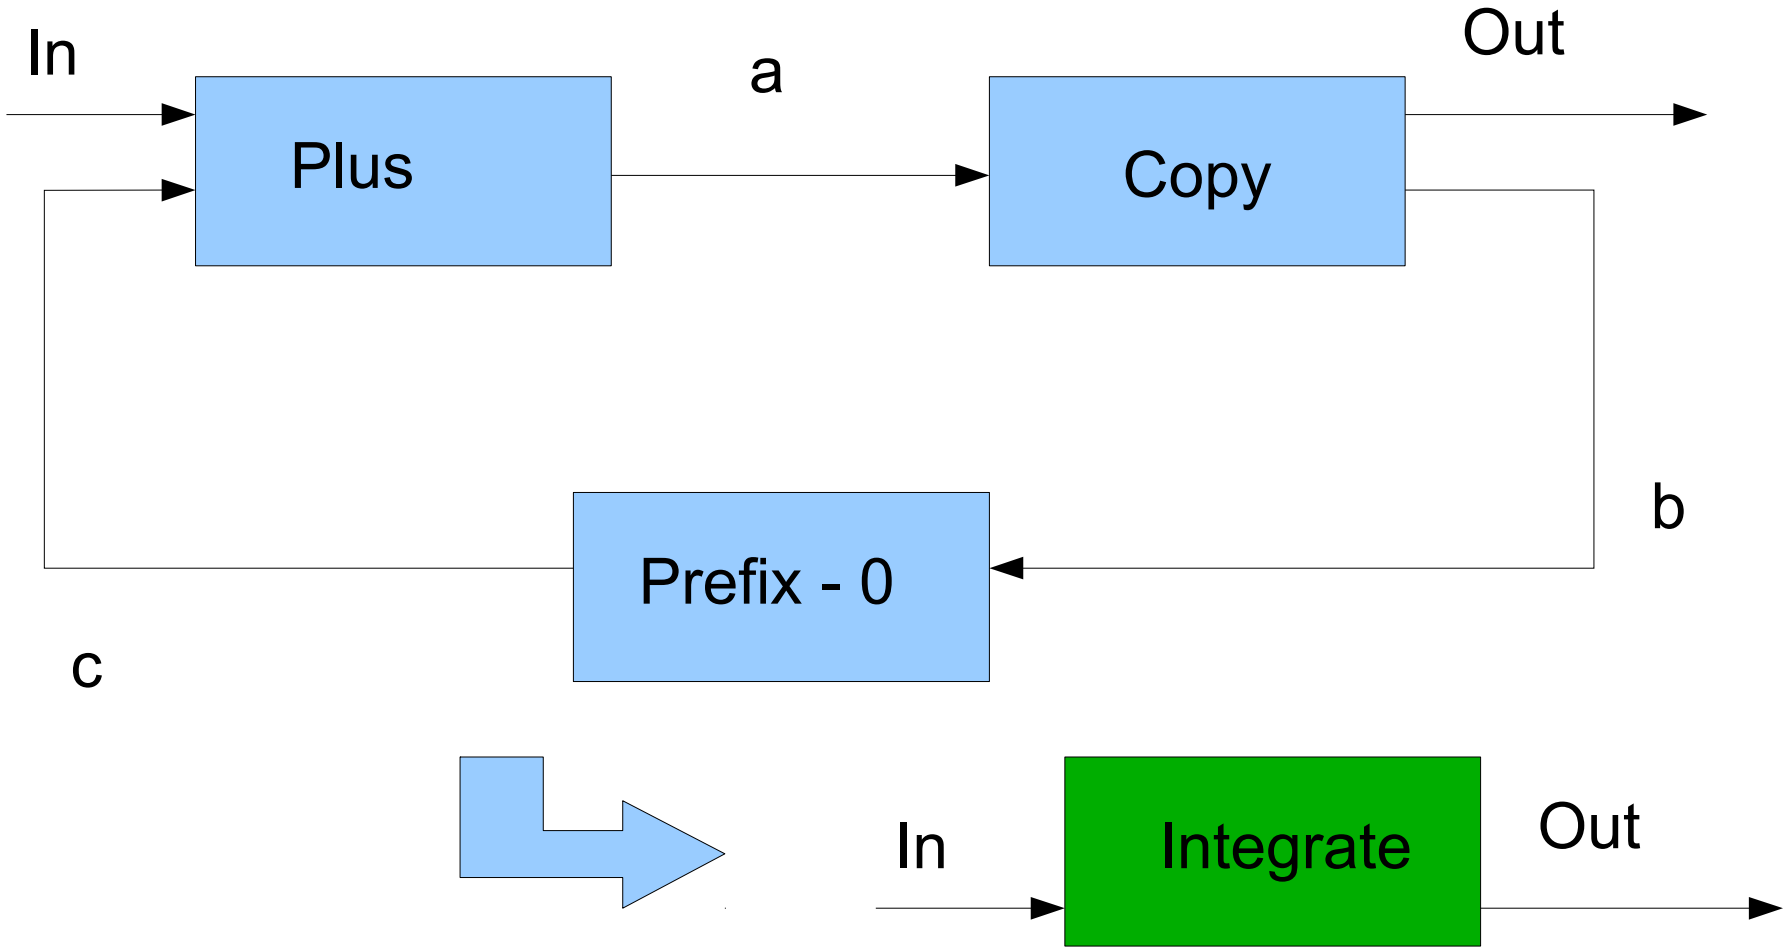

Processes

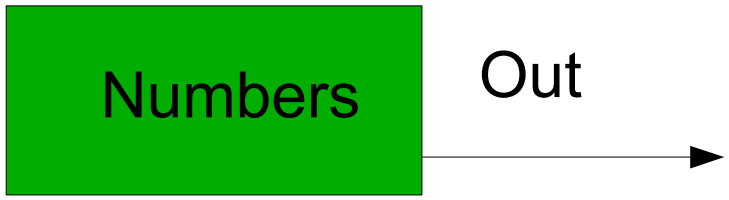

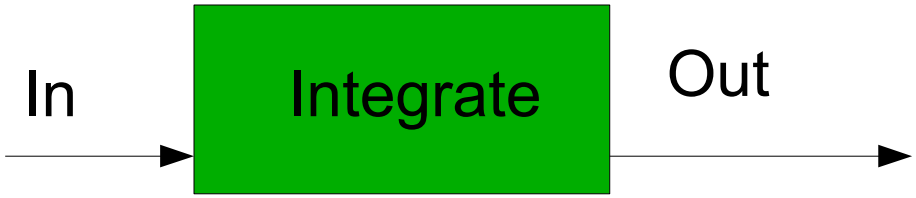

Processes

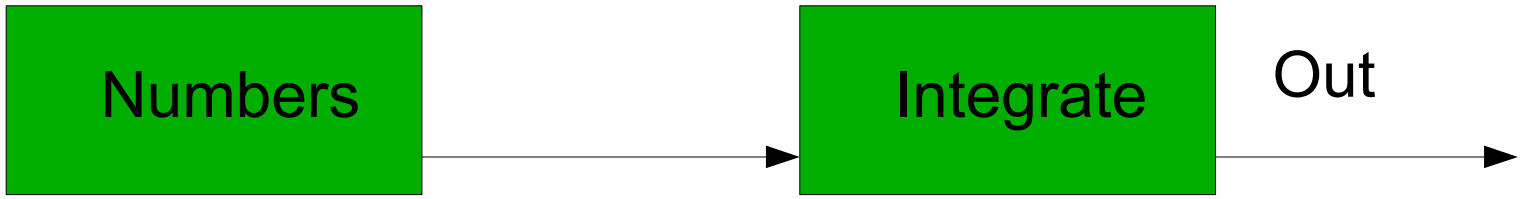

Processes

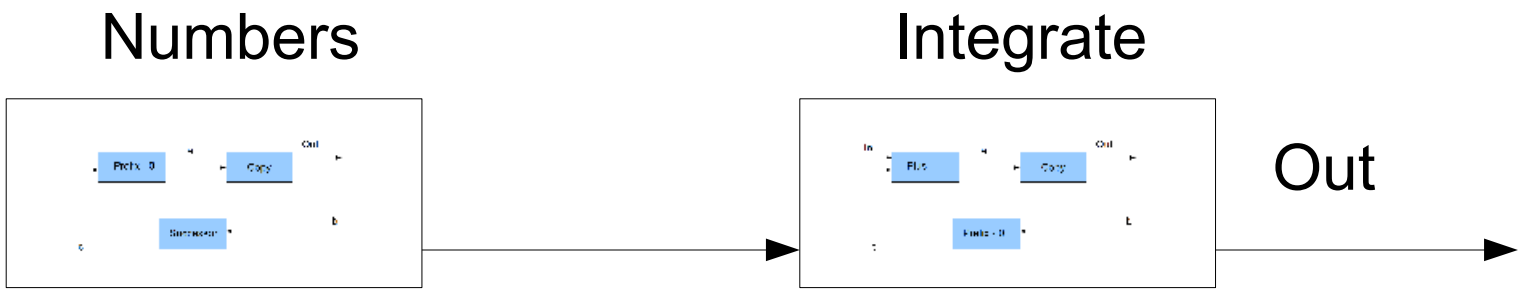

Data parallelism

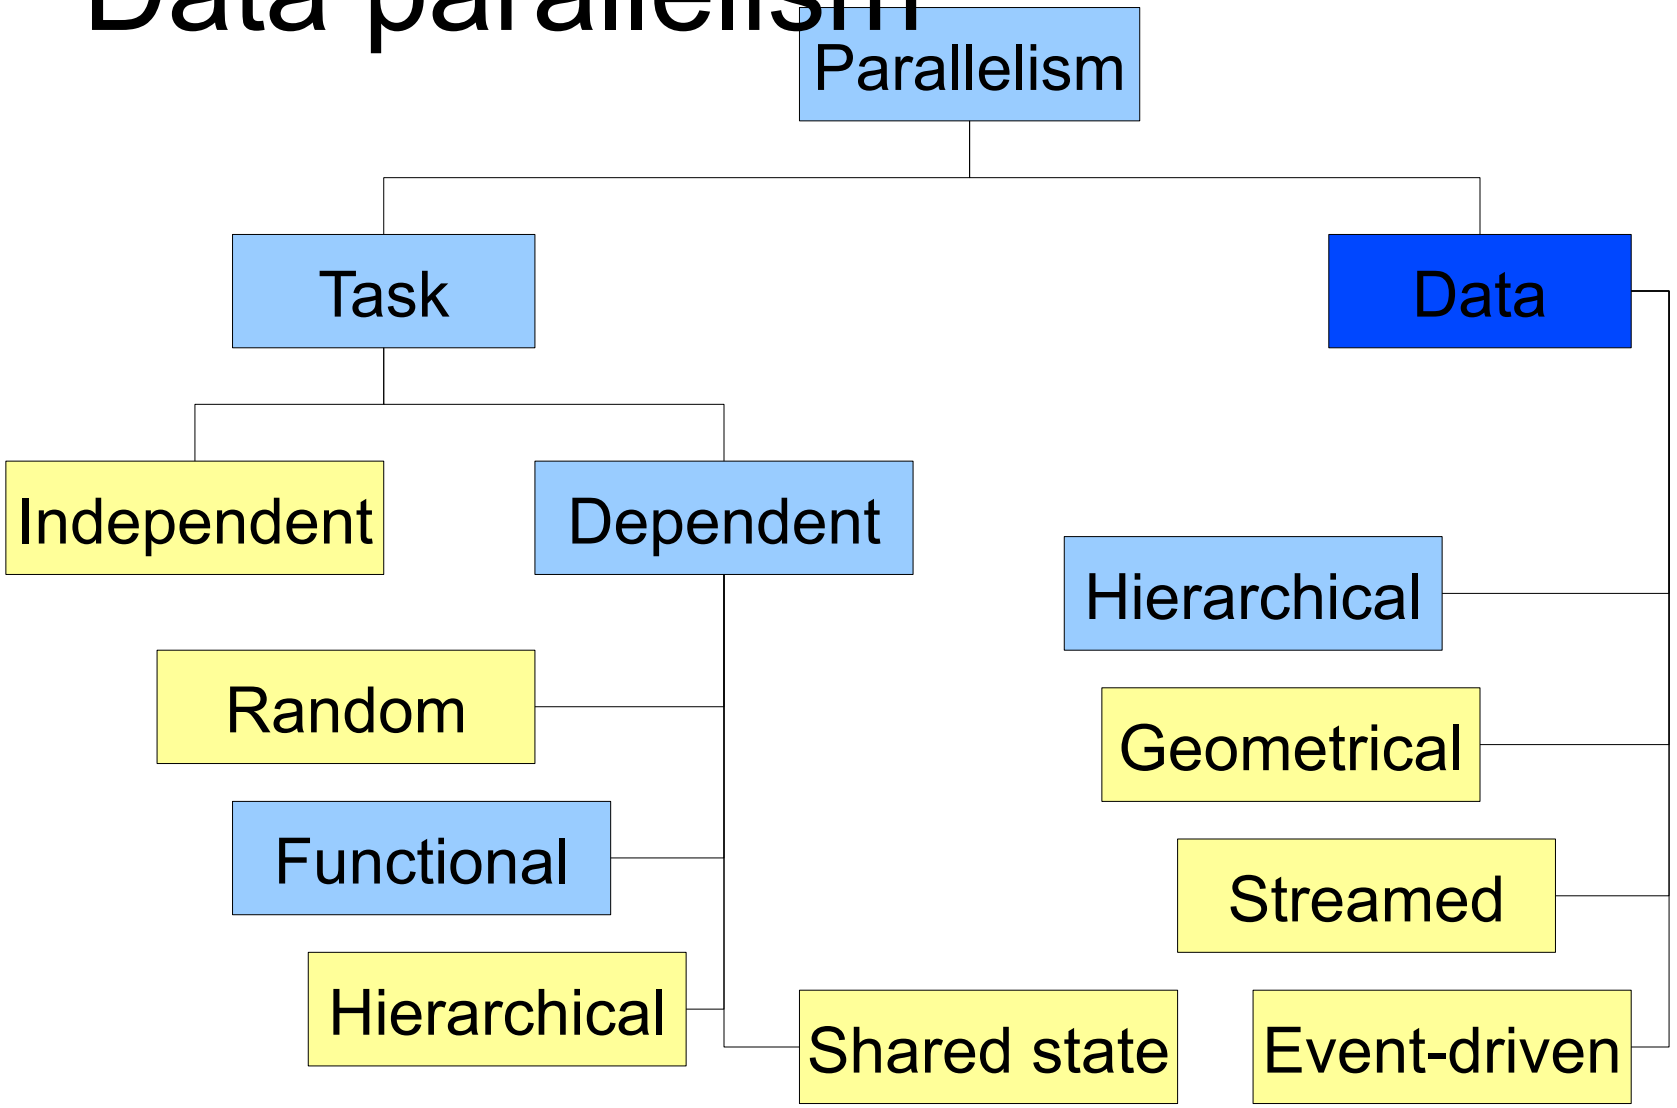

Data parallelism

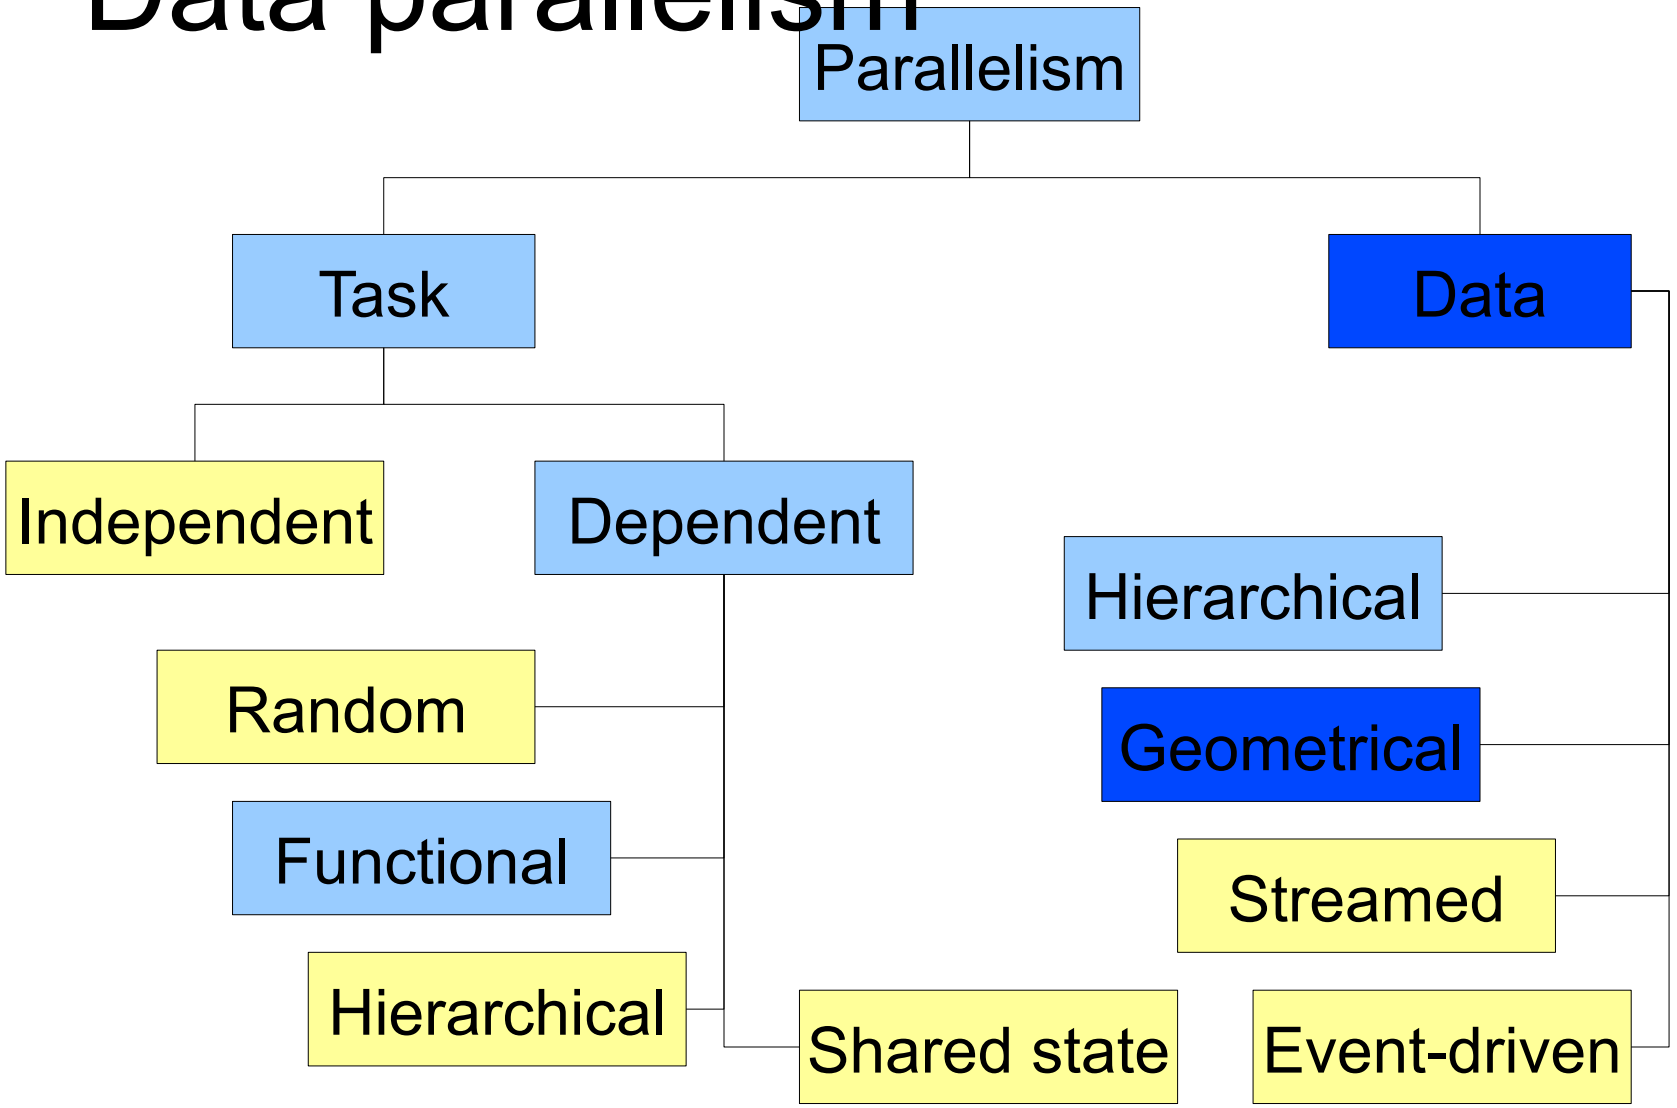

Geometric decomposition

```
images.eachParallel {it.process()}
documents.sumParallel()
candidates.maxParallel {it.salary}.marry()
```

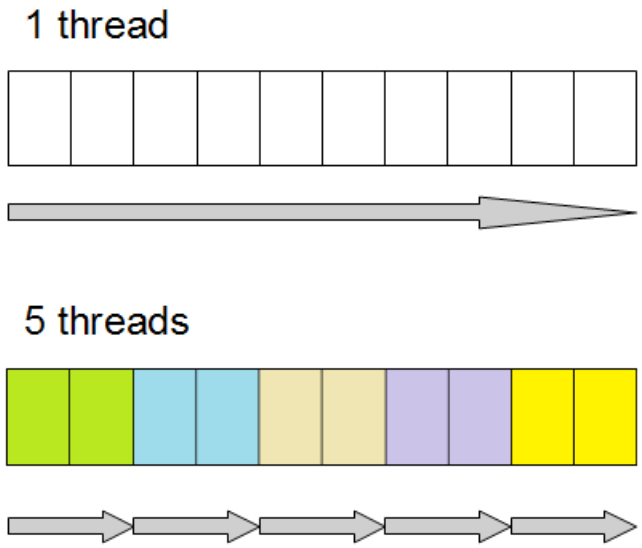

Geometric decomposition

```
registrations = submissions
        .collectParallel { form -> form.process()}
        .findAllParallel { it.valid }

registrations = submissions.parallel
        .map { form -> form.process()}
        .filter { it.valid }.collection
```

Map - reduce

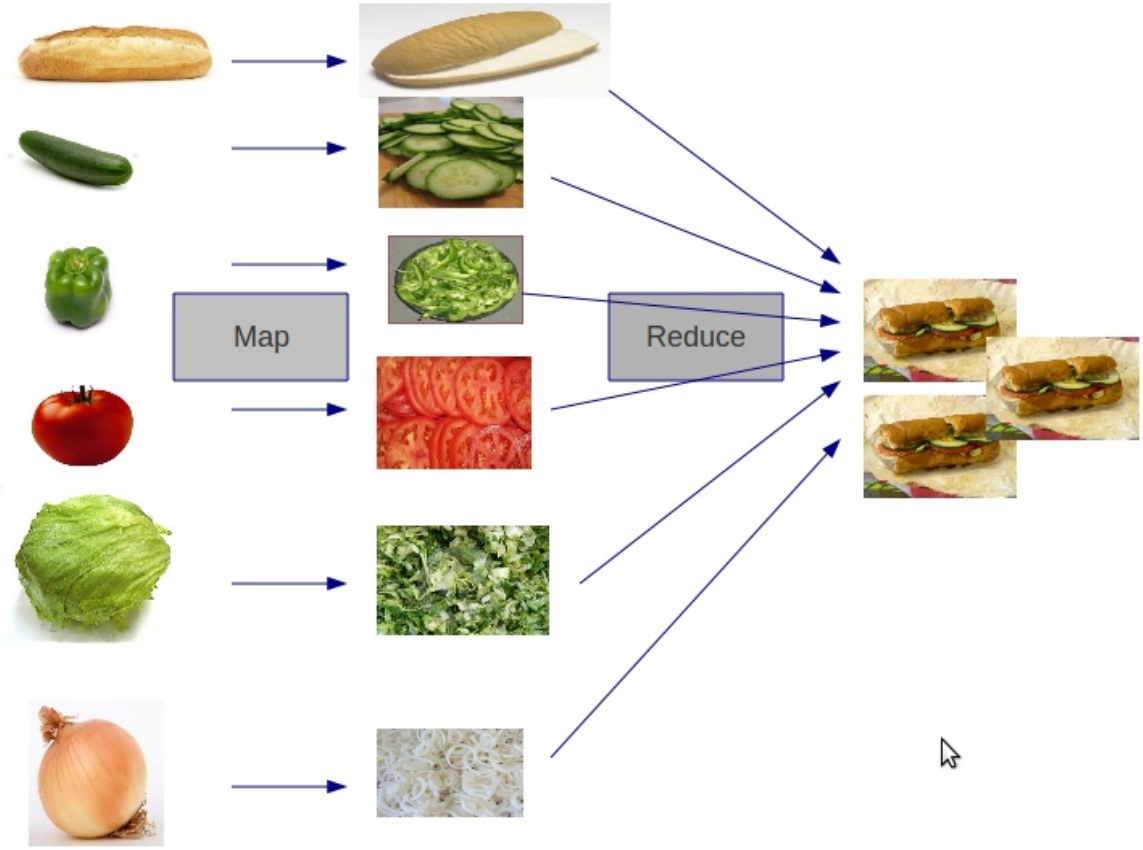

Frequent confusion

Questions | Tags | Users | Badges | Unanswered

parallel quick sort outdone by single threaded quicksort

0

I've been reading , here is the example in the book using futures to implement parallel quick sort.

But I found this function is more than twice slower than the single threaded quick sort function without using any asynchronous facilities in c++ standard library. Tested with g++ 4.8 and visual c++ 2012.

I used 10M random integers to test, and in visual c++ 2012,this function spawned 6 threads in total to perform the operation in my quad core PC.

I am really confused about the performance. Any body can tell me why?

GPU

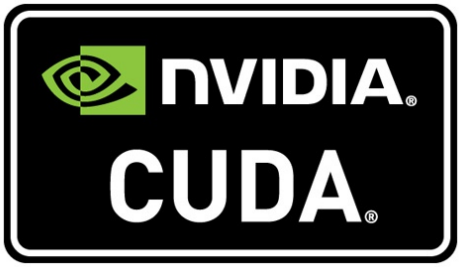

Java watch list:

https://github.com/pcpratts/rootbeer1/

http://openjdk.java.net/projects/sumatra/

Improper use 1

```
def accumulator = 0

myCollection.eachParallel {
    accumulator += calculate(it)
}
```

Improper use 2

```
new File("/file.txt").withReader{reader ->
    reader.eachParallel {
        def r1 = step1(r)
        def r2 = step2(r1)
        def r3 = step3(r2)
    }
}
```

Unroll iteration

```
def pipeline = data | step1 | step2 | step3


new File("/file.txt").withReader{reader ->
    reader.each {
        data << it
    }
}
```

Unroll iteration

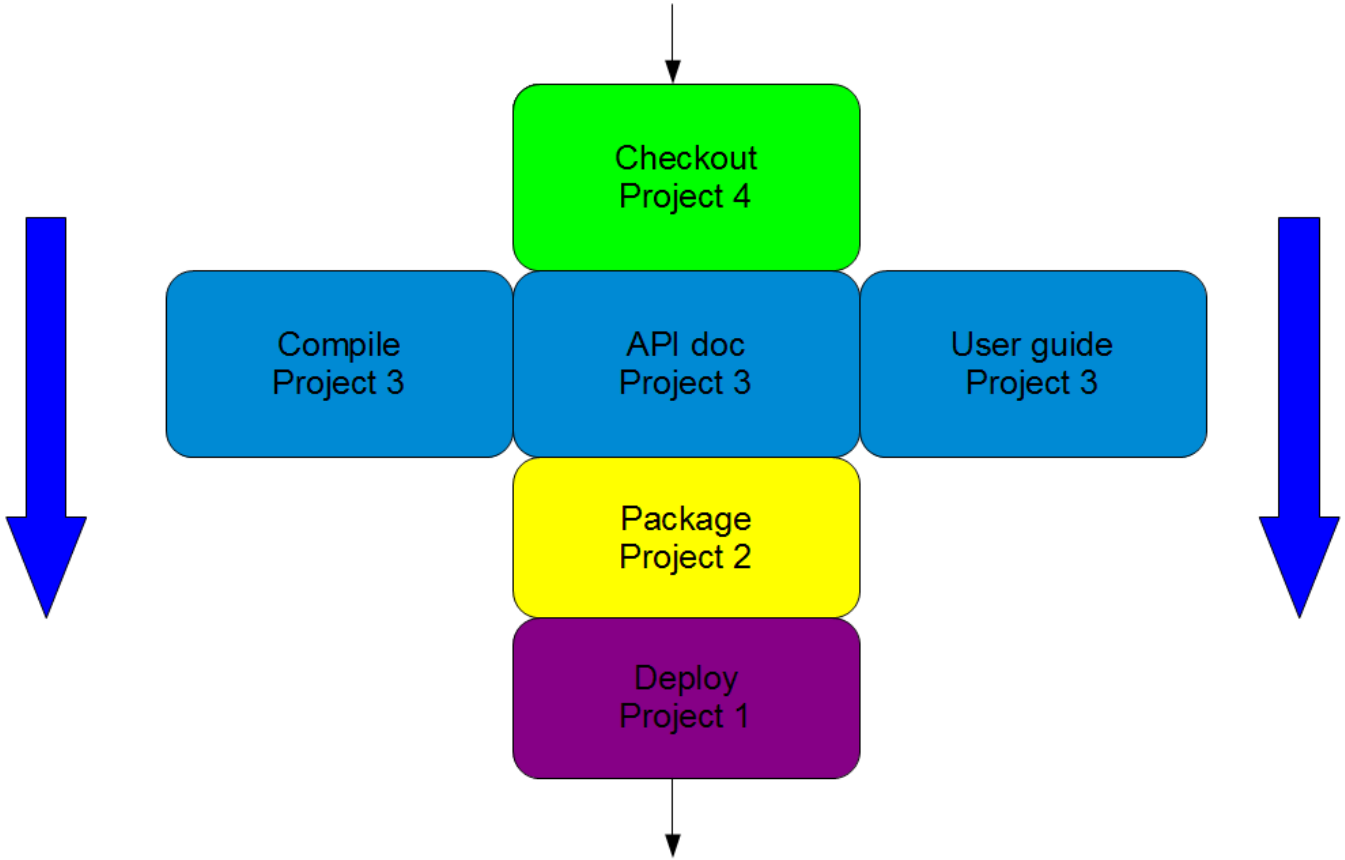

Streamed data

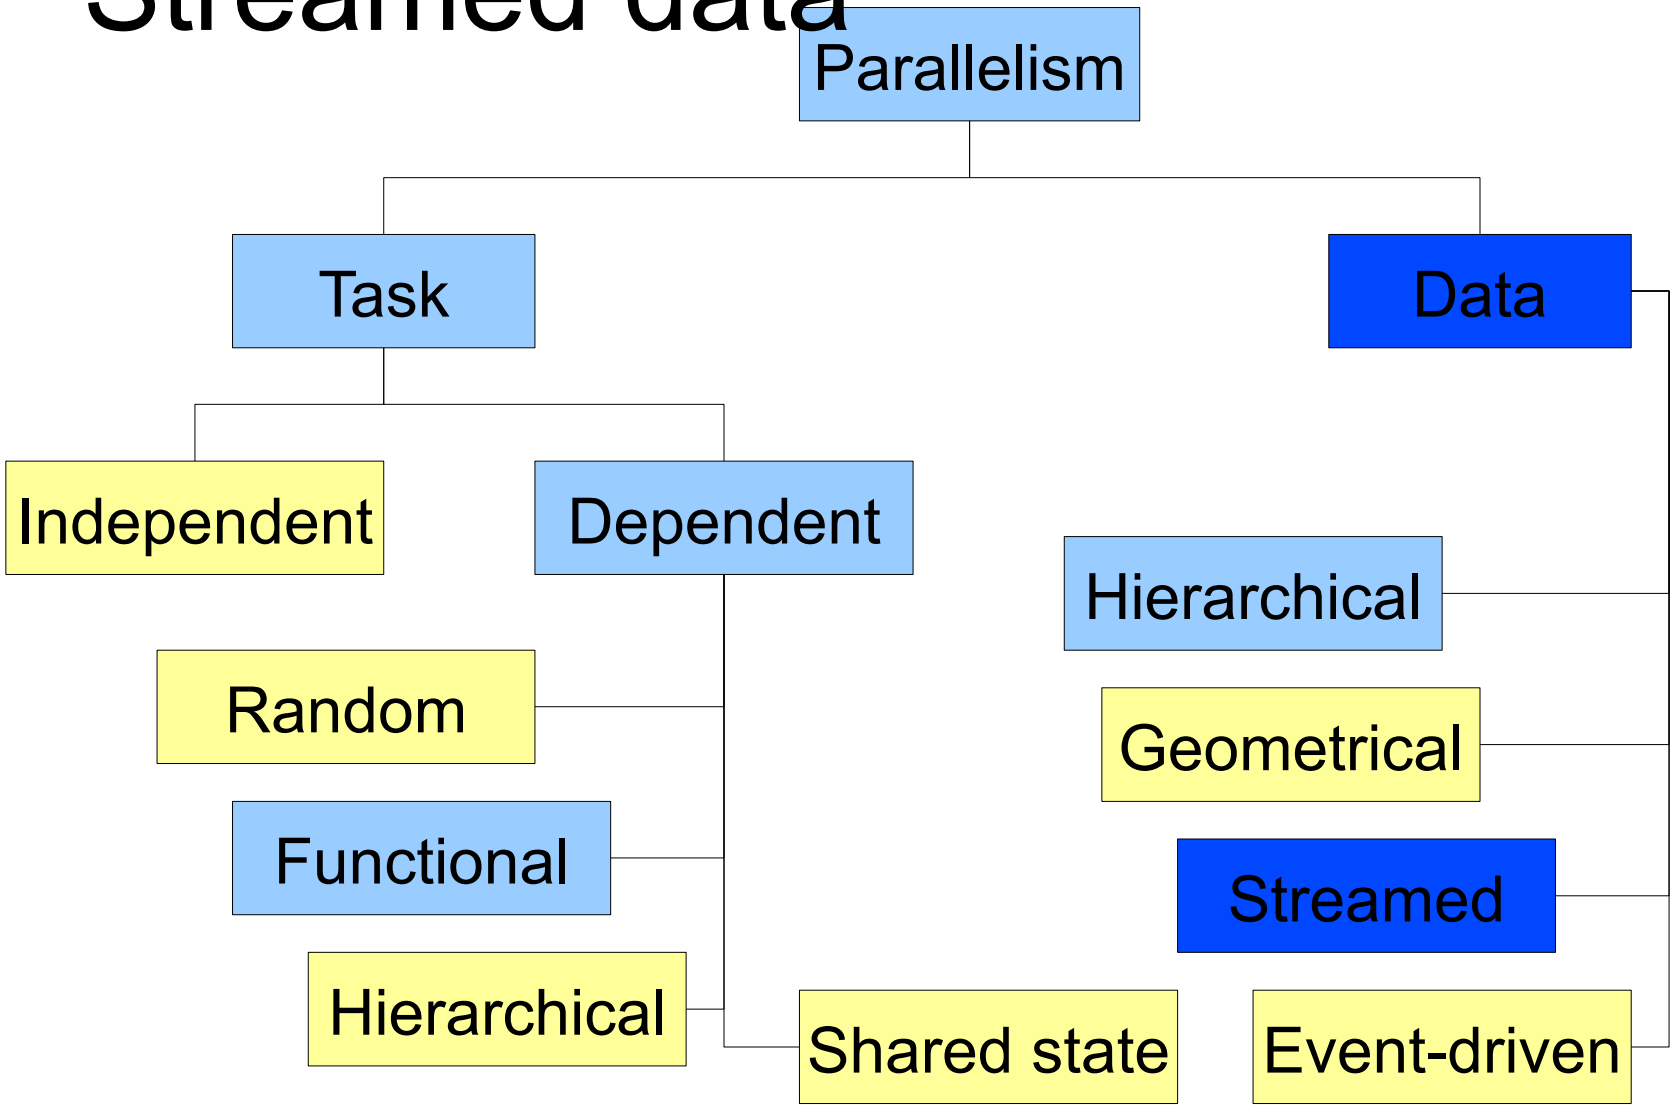

Pipeline DSL

```
def toProcess = new DataflowQueue()
def validated = new DataflowQueue()

toProcess |  {form -> process(form)} |
    {processedForm -> validate(processedForm)} | validated

submissions.each {toProcess << it}
```

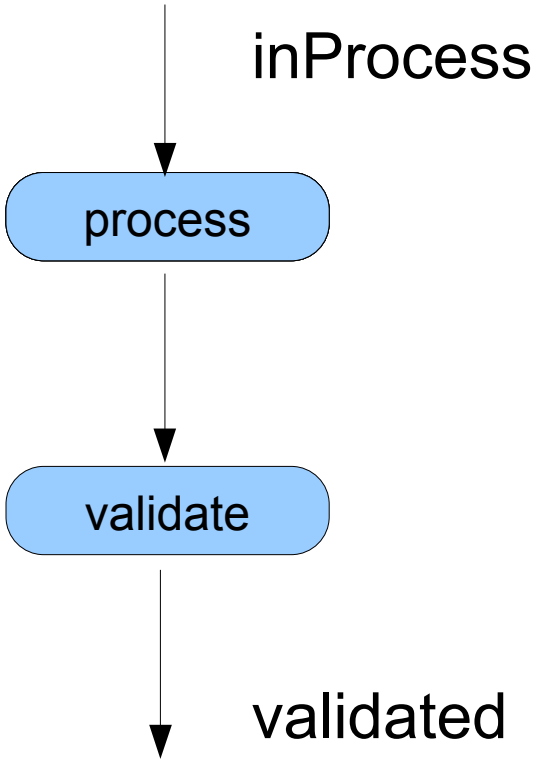
inProcess
process
validate
validated

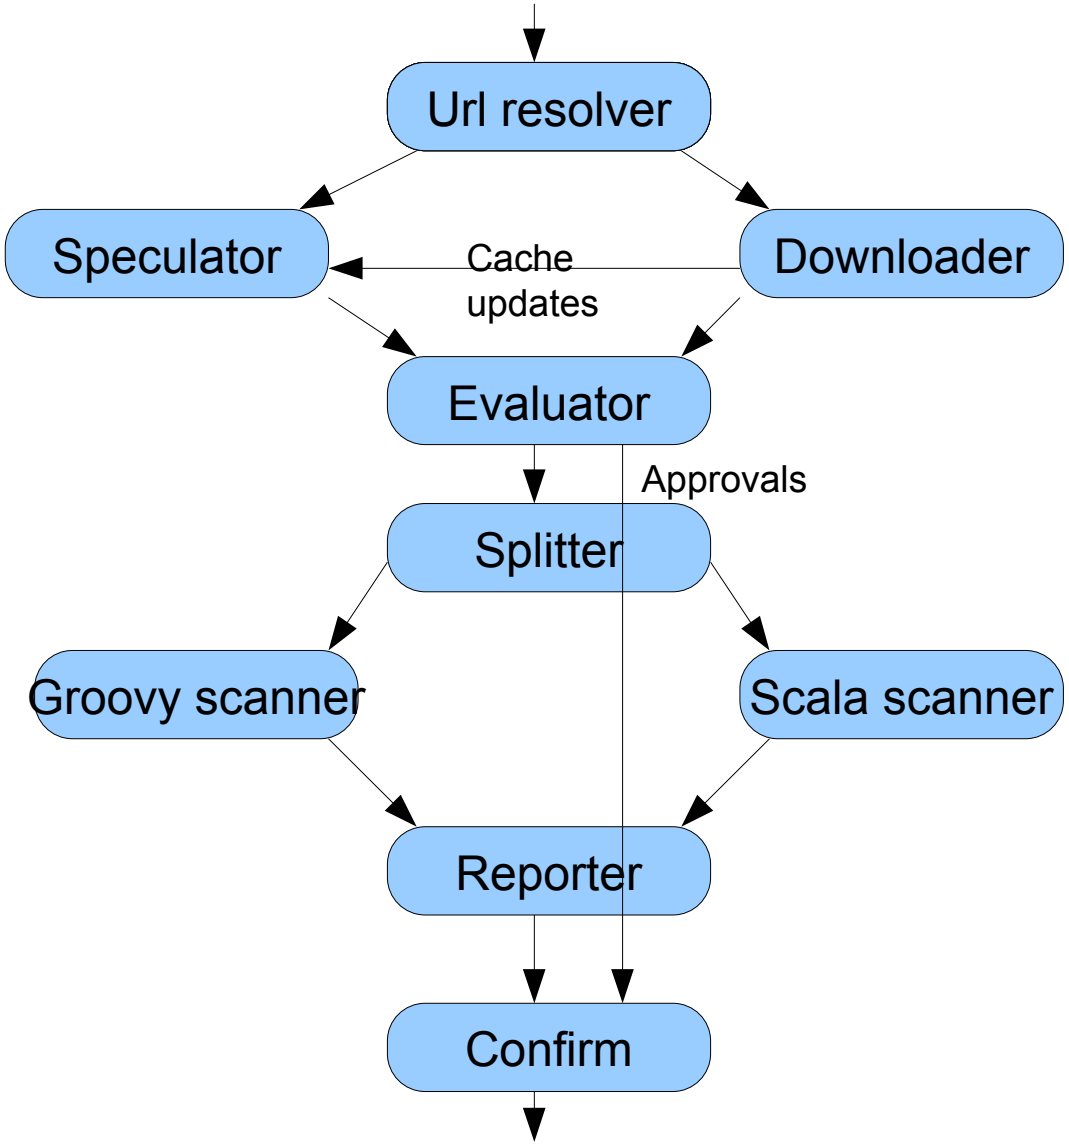
Url resolver
Speculator
Cache
updates
Downloader
Evaluator
Approvals
Splitter
Groovy scanner
Scala scanner
Reporter
Confirm

Dataflow Operators

```
operator(inputs: [headers, bodies, footers],
         outputs: [articles, summaries])
  {header, body, footer ->
    def article = buildArticle(header, body, footer)
    bindOutput(0, article)
    bindOutput(1, buildSummary(article))
  }
```

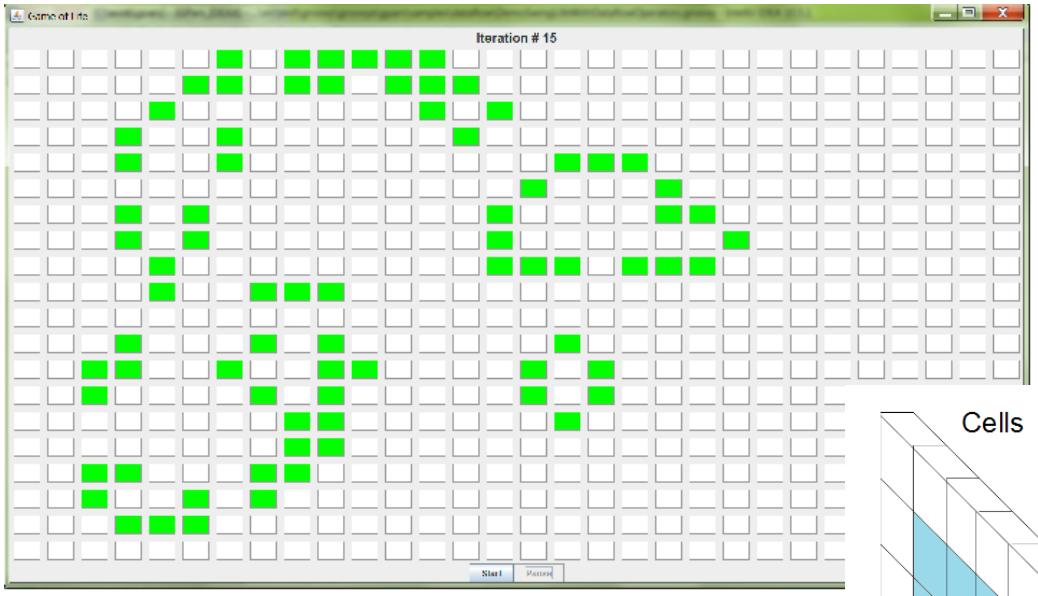
Iteration # 15
Start

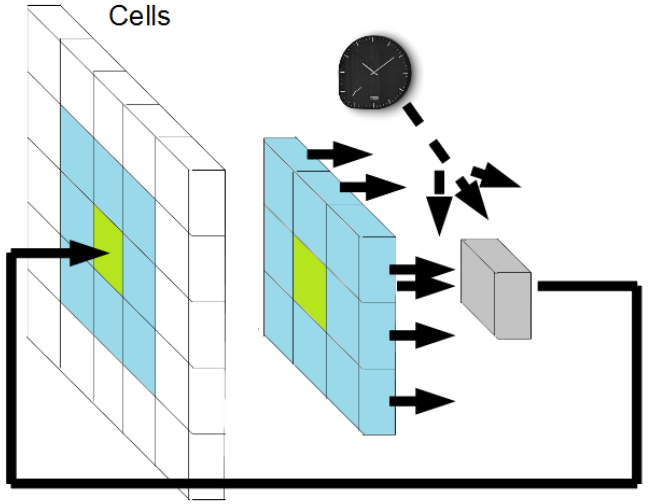
Cells

Actors

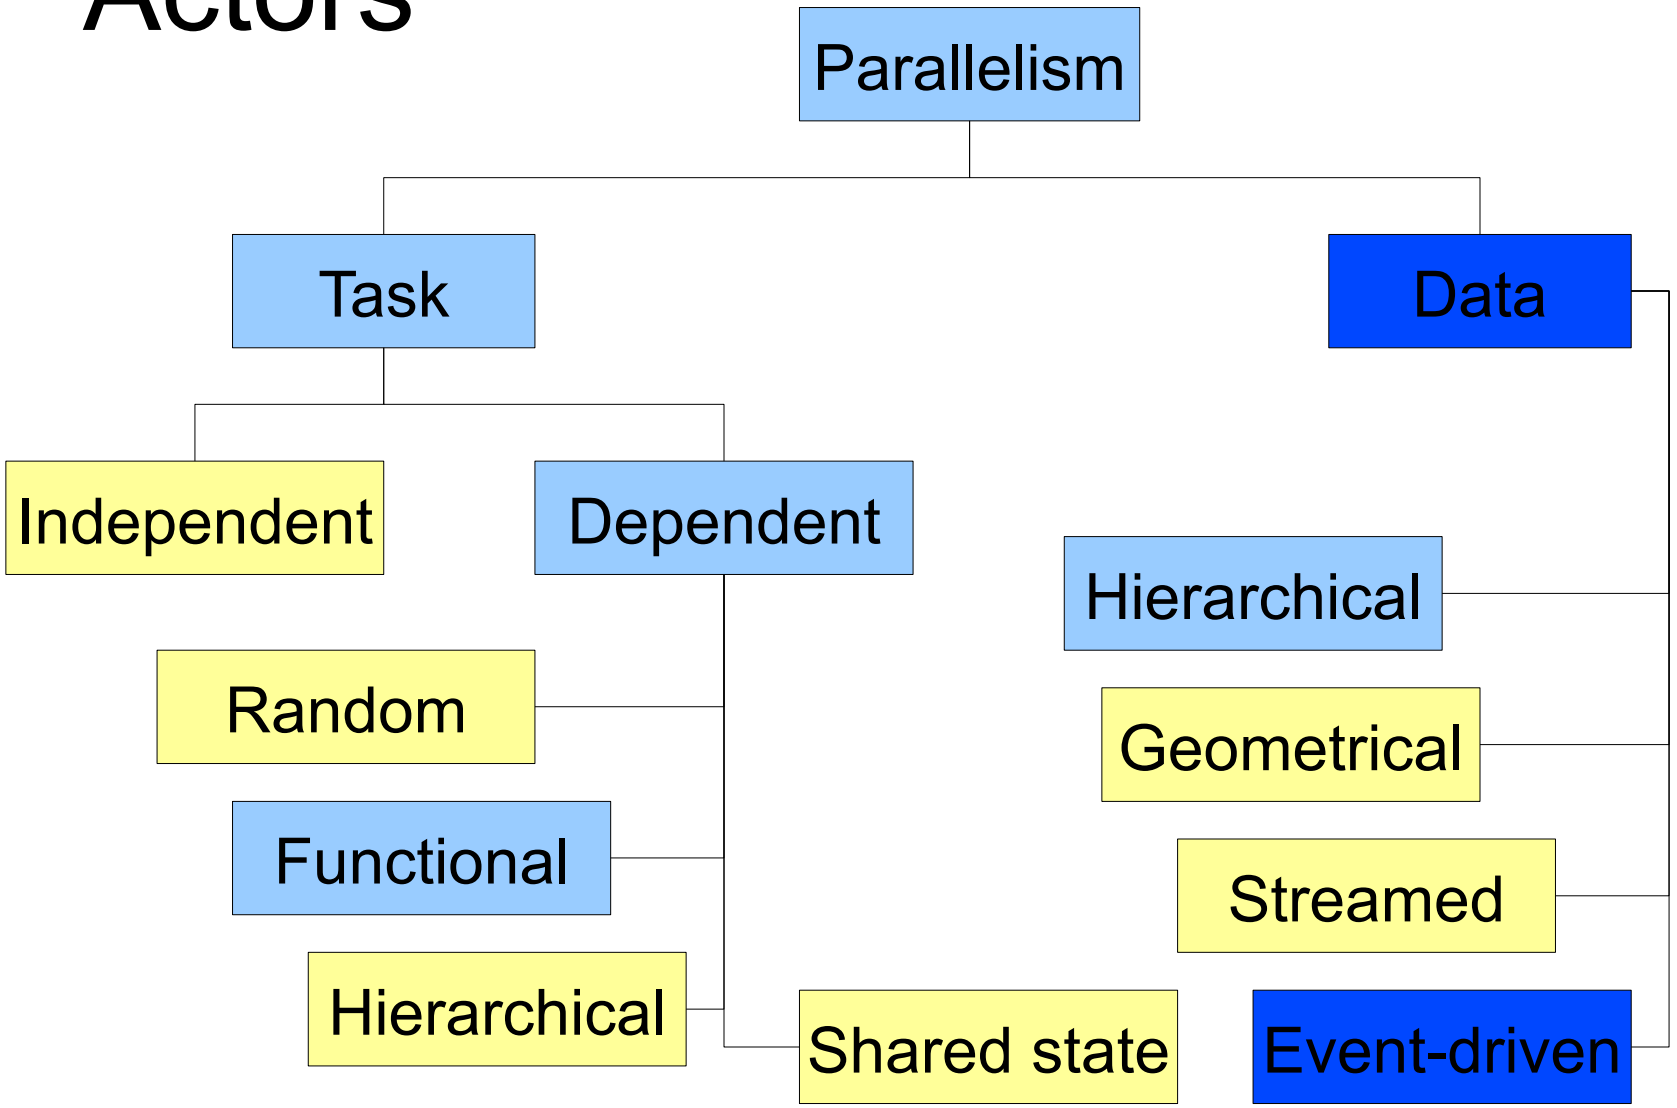
Parallelism
Task
Data
Independent
Dependent
Hierarchical
Random
Geometrical
Functional
Streamed
Hierarchical
Shared state
Event-driven

Actors

- Isolated
- Communicating
  - Immutable messages
- Active
  - Pooled shared threads
- Activities
  - Create a new actor
  - Send a message
  - Receive a message

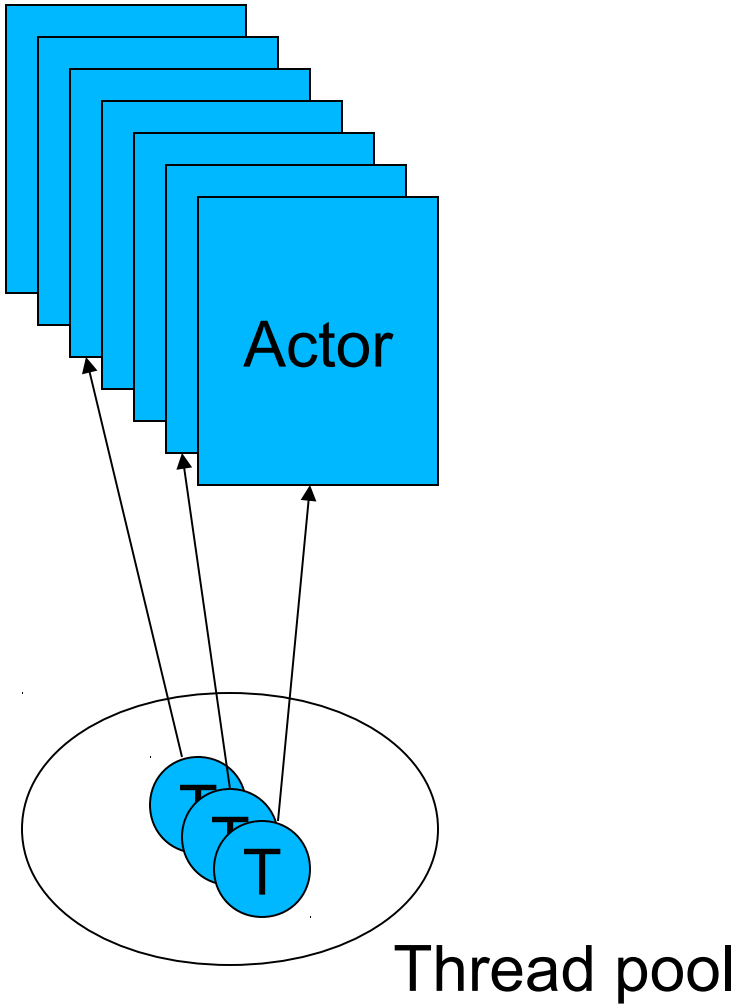

Actors use

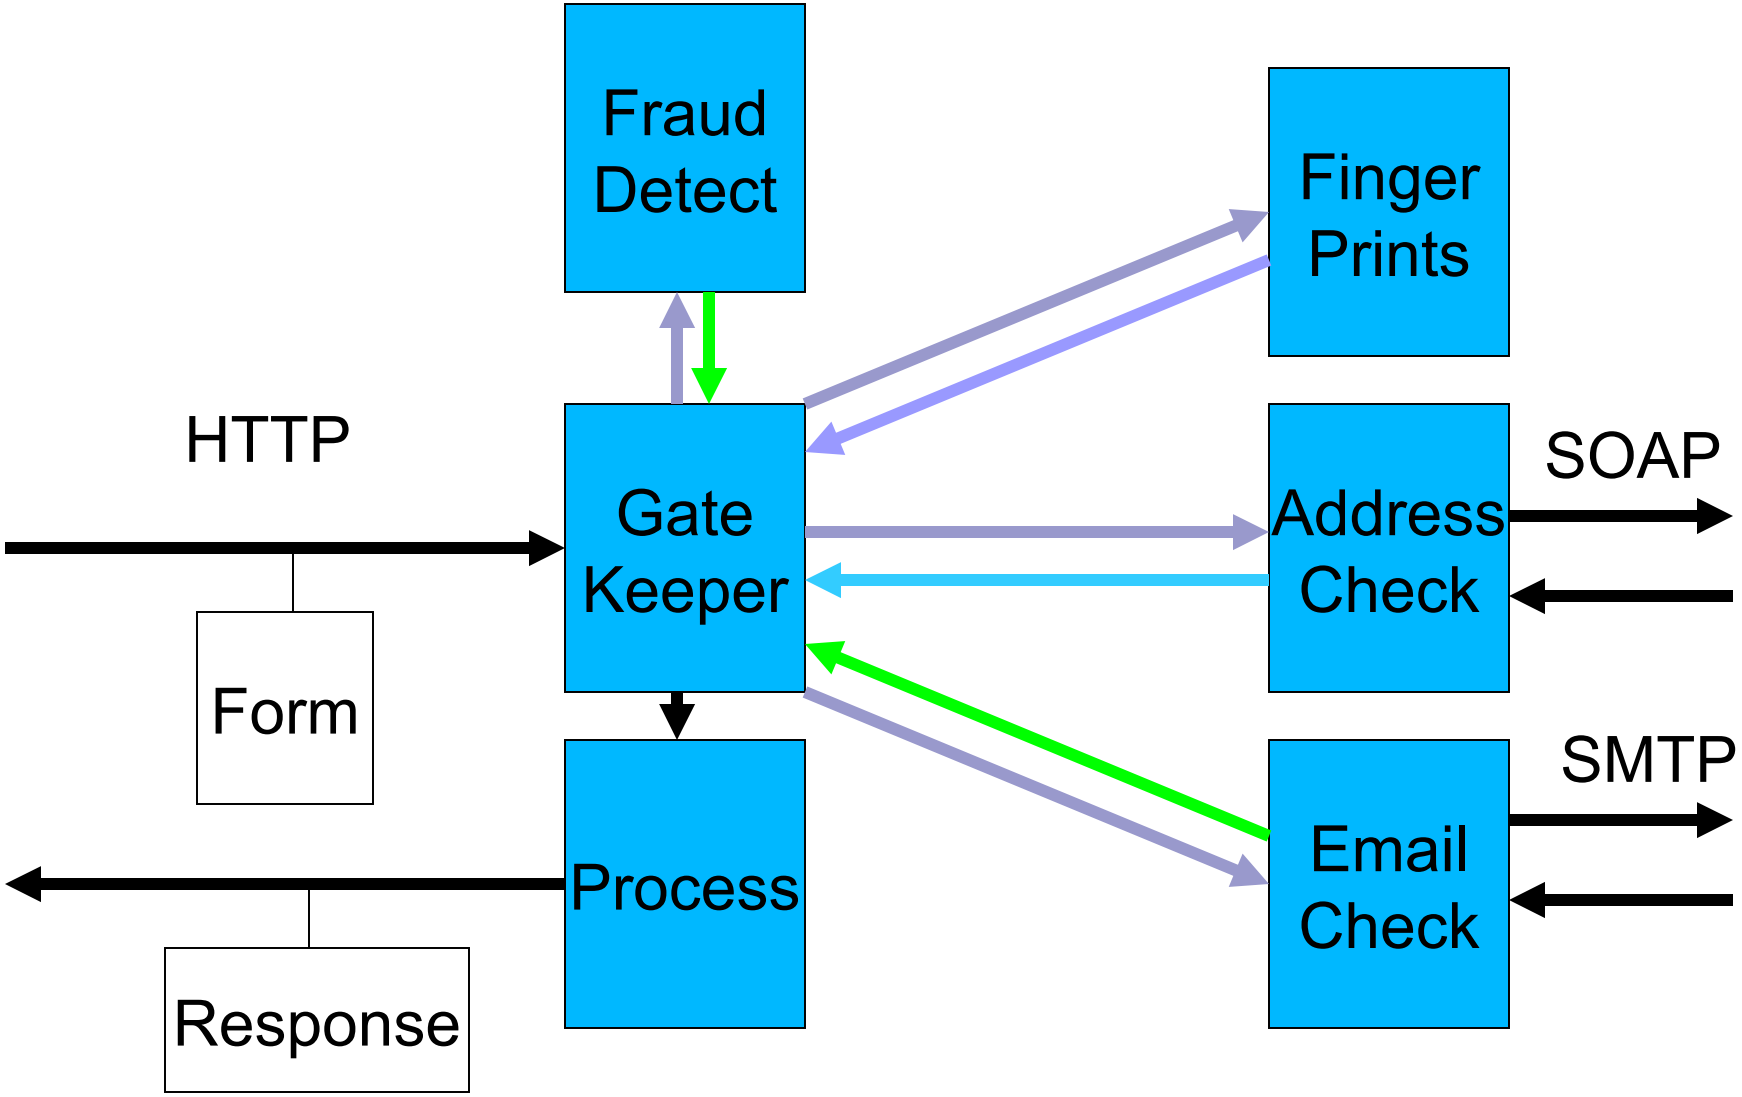
Fraud Detect
Finger Prints
HTTP
Gate Keeper
SOAP
Address Check
Form
Process
SMTP
Email Check
Response

Actors patterns

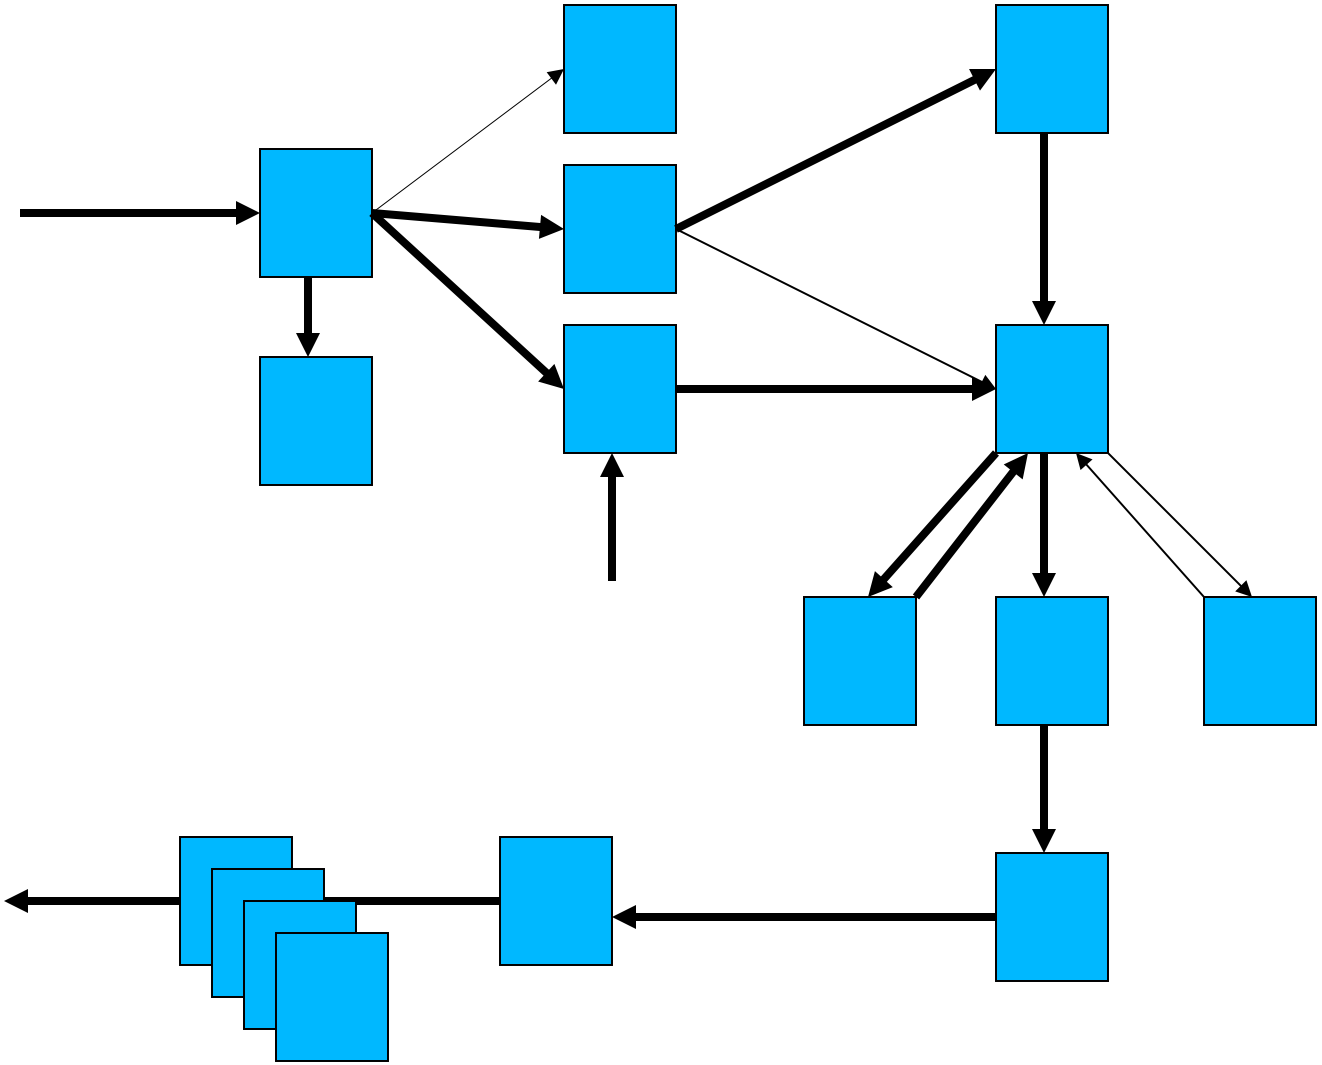

Enricher

Router

Translator

Endpoint

Splitter

Agregator

Filter

Resequencer

Checker

Sending messages

```
buddy.send 10.eur

buddy << new Book(title:'Groovy Recipes',
                  author:'Scott Davis')

def canChat = buddy.sendAndWait 'Got time?'

buddy.sendAndContinue 'Need money!', {cash->
    pocket.add cash
}
```

Event driven – actors

```
class MyActor extends DynamicDispatchActor {
    private int counter = 0

    public void onMessage(String msg) {
        this.counter += msg.size()
    }
    public void onMessage(Integer number) {
        this.counter += number
    }
    public void onMessage(Money cash) {
        this.counter += cash.amount
        reply 'Thank you'
    }
}
```

Event driven – active objects

```
@ActiveObject
class MyCounter {
    private int counter = 0


 @ActiveMethod
    def incrementBy(int value) {
        println "Received an integer: $value"
        this.counter += value
    }
}
```

Conclusion

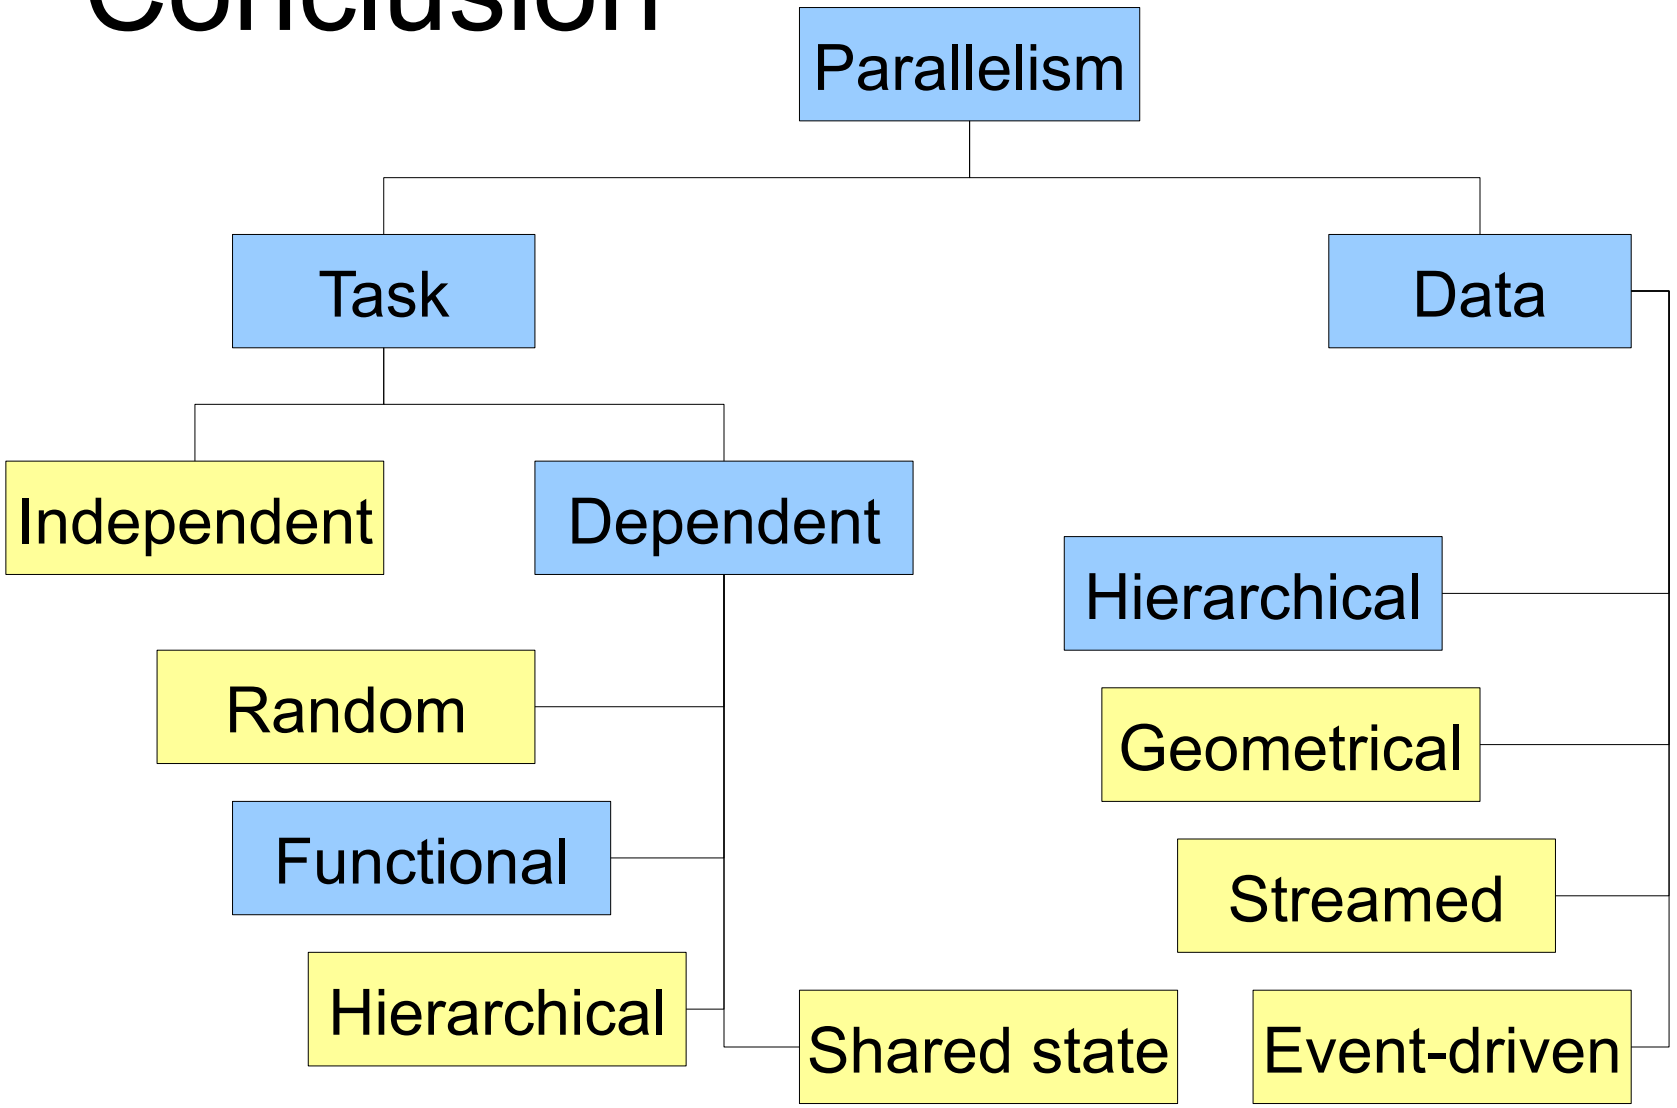
Parallelism
Task
Data
Independent
Dependent
Hierarchical
Random
Geometrical
Functional
Streamed
Hierarchical
Shared state
Event-driven

Summary

Parallelism is not hard, multi-threading is

Jon Kerridge, Napier University

Questions?

Find more at:

http://gpars.org

http://www.jroller.com/vaclav

http://twitter.com/vaclav_pech